



Greedy algorithms 2

Caching, Compression

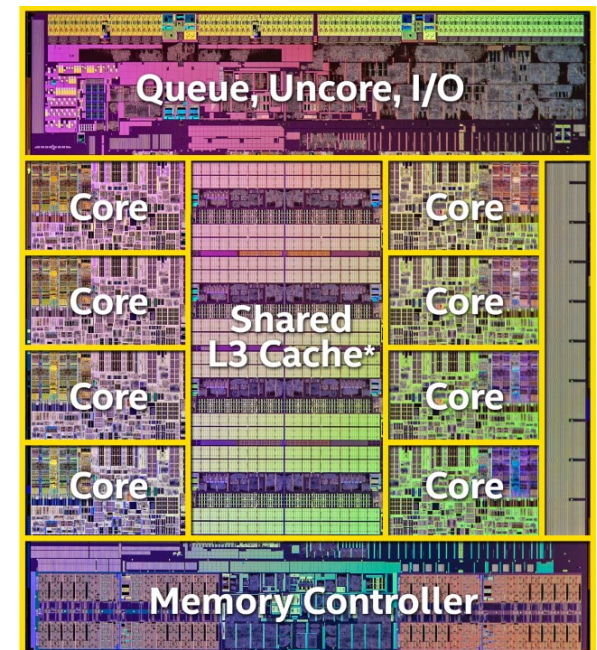

CS240

Spring 2023

Rui Fan

Caching

- Cache is a piece of on-chip (fast) memory.
- Reduces effective access time to slow main memory.
 - When accessing memory, first try to find it in cache.
 - Only if it's not in cache do we access main memory.
 - Typically cache has ~1-20 cycles latency, memory has ~200-500 cycles latency.
- Since caches are on-chip, amount is usually quite small (~32 KB to ~4 MB).
 - Only store the most important, frequently accessed data in cache.
- Use caching algorithm to select which data to store.



Optimal Offline Caching

Caching.

- Cache with capacity to store k items.
- Sequence of m item requests d_1, d_2, \dots, d_m .
- Cache hit: item already in cache when requested.
- Cache miss: item not already in cache when requested: must load requested item from main memory into cache, and evict some existing item, if full.

Goal. Eviction schedule that minimizes number of evictions.

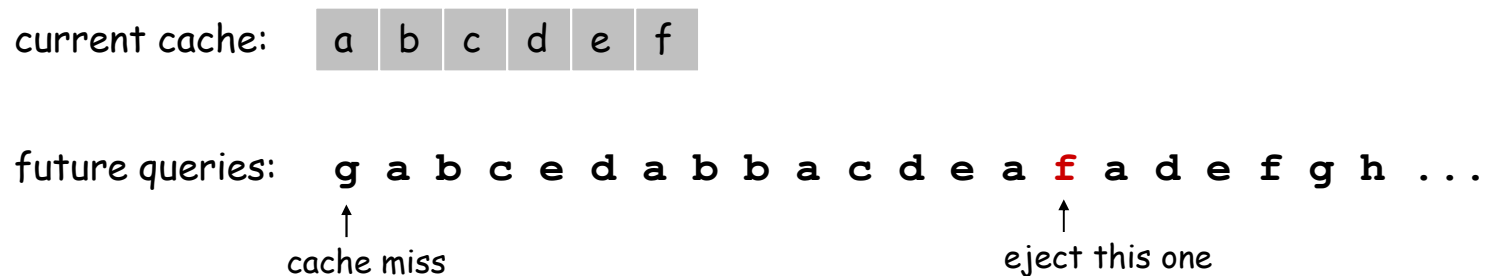
Ex: $k = 2$, initial cache = ab ,
requests: a, b, c, b, c, a, a, b .

Optimal eviction schedule: 2 evictions.

a	a	b
b	a	b
c	c	b
b	c	b
c	c	b
a	a	b
a	a	b
b	a	b
requests	cache	

Optimal Offline Caching: Farthest-In-Future

Farthest-in-future. Evict item in the cache that is not requested until farthest in the future.



Theorem. [Bellady, 1960s] FF is optimal eviction schedule.

Pf. Algorithm and theorem are intuitive; proof is subtle.

Reduced Eviction Schedules

Def. A **reduced** schedule is a schedule that only inserts an item into the cache in a step in which that item is requested.

Intuition. Can transform an unreduced schedule into a reduced one with no more evictions.

a	a	b	c
a	a	x	c
c	a	d	c
d	a	d	b
a	a	x	b
b	a	c	b
c	a	c	b
a	a	b	c
a	a	b	c

an unreduced schedule

a	a	b	c
a	a	b	c
c	a	b	c
d	a	d	c
a	a	d	c
b	a	d	b
c	a	c	b
a	a	c	b
a	a	c	b


a reduced schedule

Reduced Eviction Schedules


Claim. Given any unreduced schedule S , can transform it into a reduced schedule S' with no more evictions.

Pf. (by induction on number of unreduced items) ← doesn't enter cache at requested time

- Suppose S brings d into the cache at time t , without a request.
- Let c be the item S evicts when it brings d into the cache.
- Case 1: d evicted at time t' , before next request for d .
 - S' has one less eviction than S by time t' .
- Case 2: d requested at time t' before d is evicted.
 - S' and S have same number of evictions by time t' .



Case 1



Case 2

Farthest-In-Future: Analysis

Lemma. Let S be a reduced schedule that makes the same schedule as S_{FF} through the first j requests. Then there is a reduced schedule S' that makes the same schedule as S_{FF} through the first $j+1$ requests, and incurs no more evictions than S does.


Pf.

- Consider $(j+1)^{st}$ request $d = d_{j+1}$.
- Since S and S_{FF} have agreed up until now, they have the same cache contents before request $j+1$.
- Case 1: (d is already in the cache). $S' = S$
- Case 2: (d is not in the cache and S and S_{FF} evict the same element).
 $S' = S$

Farthest-In-Future: Analysis

Pf. (continued)

- Case 3: (d is not in the cache; S_{FF} evicts e ; S evicts $f \neq e$).
 - begin construction of S' from S by evicting e instead of f



- now S' agrees with S_{FF} on first $j+1$ requests
- From request $j+2$ onward, we make S' the same as S , but this becomes impossible when e or f is involved

Farthest-In-Future: Analysis

Pf. (continued)

Let j' be the **first** time after $j+1$ that S and S' take a different action, and let g be item requested at time j' .




- Case 3a: $g = e$.
 - f couldn't have been requested between time $j+1$ and j' , because if it had been, S and S' would have taken different actions before j' .
 - So e is requested before f . But this can't happen with Farthest-In-Future, since S_{FF} evicted e , implying f is requested before e .

Farthest-In-Future: Analysis

Pf. (continued)

Let j' be the **first** time after $j+1$ that S and S' take a different action, and let g be item requested at time j' .




S can't evict f , and if it evicts $e' \neq e, f$, then S' , by construction, would do the same thing.



- Case 3b: $g \neq e, f$. S must evict e .


Make S' evict f ; now S and S' have the same cache. ▪



Farthest-In-Future: Analysis

Pf. (continued)

Let j' be the **first** time after $j+1$ that S and S' take a different action, and let g be item requested at time j' .



- Case 3c: $g = f$. Element f can't be in cache of S , so let e' be the element that S evicts.
 - if $e' = e$, S' accesses f from cache; now S and S' have same cache
 - if $e' \neq e$, S' evicts e' and brings e into the cache; now S and S' have the same cache.
 - S' is no longer reduced, but can be transformed using procedure on slide 6 into a reduced schedule which
 - a) agrees with S_{FF} through step $j+1$
 - b) has no more evictions than S

Farthest-In-Future: Analysis

Theorem. FF is optimal eviction algorithm.

Pf. (by induction on number of requests j)

Base case (trivial):

There exists an optimal reduced schedule S that makes the same schedule as S_{FF} through the first 0 requests.

Inductive step (implied by the lemma):

If there exists an optimal reduced schedule S that agrees with S_{FF} through the first j requests,

then there exists an optimal reduced schedule S' that agrees with S_{FF} through the first $j+1$ requests

Caching Perspective

Online vs. offline algorithms.

- Offline: full sequence of requests is known a priori.
- Online (reality): requests are not known in advance.
- Caching is among most fundamental online problems in CS.

LIFO. Evict page brought in most recently.

LRU. Evict page whose most recent access was earliest.



↑
FF with direction of time reversed!

Theorem. FF is optimal offline eviction algorithm.

- Provides basis for understanding and analyzing online algorithms.
- LRU is k -competitive.
 - I.e. it does $\leq k$ times more loads than the optimal eviction algorithm.
- LIFO is arbitrarily bad.
 - For n requests, it may do $O(n)$ times more loads than optimal.

Compression

- Storing and transmitting data is expensive. Compression represents data more compactly.
- ASCII has 256 characters, so we use $\log_2(256)=8$ bits to represent each character.
- But typically some characters **appear more often** than others. So we shouldn't use same number of bits for all letters.
- Basic idea for compression is to use different length bitstrings.
 - Use **short bitstrings** to represent **common characters**.
 - Use **long bitstrings** to represent **uncommon characters**.
 - We save space on average.





Lossless vs. lossy compression

- Different algorithms for different applications.
- **Lossless** compression used in settings where losing even one bit can make data useless.
 - **Ex** Computer code, financial document.
 - Typical compression ratio is 2:1.
 - Huffman encoding is lossless.
- **Lossy** compression used when data still useful after losing some information.
 - **Ex** Audio and video, MP3 and MPEG.
 - We can't hear high frequencies or see fast movement, so we can discard this info.
 - Typical compression ratio is 5:1-50:1.



Variable length encoding

- Let's compress “**lollapalooza**”.
- There are 5 different letters. If we use the same length bitstring to represent each letter, we need 3 bits per letter.
 - We use 36 bits total.
- To use different length bitstrings, first count how many times each letter appears.
 - 4 l's, 3 a's, 3 o's, 1 p, 1z.
- Use shorter bitstrings for more frequent letters.
- Use the encoding **l=00, a=01, o=10, p=110, z=111**.
 - Encoding of “lollapalooza” is **00100000011100100101011101**, formed by replacing each letter by its encoding.
 - We use 26 bits, for a 28% savings.



Ambiguity

- We want the codewords to be short, but we also need them to be **unambiguously decodable**.
- **Ex** If we use $l=0$, $a=01$, $o=10$, $p=110$, $z=111$, then “lollapalooza” is only 22 bits.
 - But we can’t decode this encoding!
 - If we see 00010, we can’t tell whether this is encoding $lla=[0,0,01,0]$, or $llo=[0,0,0,10]$.
- We could use a separator, $0\#0\#01\#0$ vs $0\#0\#0\#10$. But that’s wasteful.
- Instead, we use **prefix-free codes**, which are unambiguously decodable.



Prefix-free codes

- Let W be the set of codewords we use. Then W is prefix-free if no codeword in W is a prefix of another codeword.
 - 00, 10, 001, 100 is not prefix-free.
 - 00 is a prefix of 001, and 10 is a prefix of 100.
 - 00, 01, 10, 110, 111 is prefix-free.
- Prefix-free codes allow **unique decoding**.
 - Given the encoded string, just keep reading until you've read a complete codeword.
 - This codeword can't be part of a longer codeword, because the code is prefix-free.

Decoding prefix-free codes


- Let $S=00100000011100100101011101$,
 $W=\{00,01,10,110,111\}$, representing
 l,a,o,p,z .

00100000011100100101011101	00→l
00100000011100100101011101	10→o
00100000011100100101011101	00→l
00100000011100100101011101	00→l
00100000011100100101011101	00→a
00100000011100100101011101	110→p

...

Huffman coding

- Huffman encoding is an optimal prefix-free code, invented in 1951.
- First, find the frequencies of the letters in your text.
 - For “lollapalooza”, it’s [l,a,o,p,z] → [4,3,3,1,1].
- Now, build a binary tree on the letters bottom up.
 - Make each letter a leaf, and set its weight to its frequency.
 - Take the **two lowest weight** nodes
 - Make them the children of a parent node.
 - Set the weight of the parent node equal to the sum of the weights of the two children.
 - Remove the two nodes.
 - Notice this is a **greedy** step.
 - Repeat till all nodes part of one tree.
- Represent left by 0, right by 1.
 - A letter’s encoding is represented by its **path from the root**.





Huffman coding example


l	a	o	p	z
4	3	3	1	1

Huffman coding example


l
4

a
3


o
3




Huffman coding example




Huffman coding example



Huffman coding example



Huffman coding example



Convert the Huffman tree into a binary encoding, by treating left (right) branches as 0's (1's).
The encoding of the letters corresponds to their path from the root.



Huffman implementation

Let $S=s_1s_2\dots s_n$ be a string. Let $f(s)$ be the number of occurrences of char s in S .

for $i=1$ to n

 add $(s_i, f(s_i))$ to a min-heap H

for $i=1$ to $n-1$

 left \leftarrow removeMin(H)

 right \leftarrow removeMin(H)

 make a new node parent

 set parent's left and right children to left, right

 add(parent, $f(\text{left}) + f(\text{right})$) to H

a letter's encoding is represented by its path

Huffman's complexity

Total time is $O(n^2)$.

for $i=1$ to n

add $(s_i, f(s_i))$ to a min-heap H

Each add takes $O(\log n)$ time.

for $i=1$ to $n-1$

left \leftarrow removeMin(H)

Each remove takes $O(\log n)$ time.

right \leftarrow removeMin(H)

make a new node parent

set parent's left and right children to left, right


add(parent, $f(\text{left}) + f(\text{right})$) to H

There are n letter, each has $O(n)$ height, so the total time is $O(n^2)$.

a letter's encoding is represented by its path.


Huffman code is prefix-free

- Any two codewords correspond to paths from the root to leaves.
 - The 2 paths split from each other somewhere.
 - After the split, neither codeword is a prefix of the other.
- Huffman codes are uniquely decodable.



Huffman code is optimal

- Huffman encoding gives the **shortest uniform encoding** of strings.
 - Uniform basically means you can't change your encoding method for different strings.
- Call an encoding in which codewords are derived from paths in trees a **tree code**.
- **Fact** There exists tree codes that are optimal.
- Since the Huffman code is a tree code, to prove Huffman is optimal, we just need to prove it's an optimal tree code.




Huffman code is optimal

- **Def** The cost of a tree code is $cost(T) = \sum_{v \in T} d(v) \cdot f(v)$, where $d(v)$ denotes the depth of letter v , and $f(v)$ denotes its frequency.

□ $cost(T)$ = number of bits to represent original string.


- **Claim 1** In an optimal tree code, every leaf has a sibling.
- **Proof** Otherwise, replace the lone leaf by its parent to get a tree code with lower cost.



$$cost(T) = 4 \cdot 2 + 3 \cdot 2 + 3 \cdot 2 + 3 \cdot 1 + 3 \cdot 1 = 26$$

Huffman code is optimal

- **Claim 2** Consider the two least frequent letters x and y . In the an optimal tree code T , x and y are **siblings** of each other at the **max depth** of the tree.
- **Proof** Suppose not, and let p be a node at the max depth.
 - p has a sibling q by Claim 1.
 - p and q have higher frequency than x and y , resp.
 - Create a new tree where we swap p and x , and q and y .
 - The new tree has strictly lower cost than T , since p, q have higher frequency than x, y . Contradiction.



Huffman code is optimal

- **Thm** Huffman code is optimal.
- **Proof** Use induction on number of letters in the code. Suppose it's true up to $n - 1$.
 - Consider a code with n letters. Let x, y be the letters with the lowest frequency.
 - Let T be the Huffman code tree on the n letters.
 - Create a new node z with frequency $f(z) = f(x) + f(y)$.
 - Let S be a tree formed from T by removing x and y , and replacing their parent by z .
 - S is the Huffman code on the $n - 1$ letters.
 - Because of the recursive way Huffman encoding works.
 - S is an optimal tree code on the $n - 1$ letters, by induction.
 - $cost(T) = cost(S) + f(x) + f(y)$.
 - All the nodes in S and T are the same, except x, y, z .
 - $cost(z) = d(z) \cdot f(z) = d(z) \cdot (f(x) + f(y))$.
 - $d(z) = d(x) - 1 = d(y) - 1$.
 - $cost(x) + cost(y) = cost(z) + f(x) + f(y)$.

Huffman code is optimal

■ Proof (continued)

- Let T' be an optimal tree code on the n letters.
 - By Claim 2, x and y are siblings in T' .
 - Merge them into a node z' , with $f(z') = f(x) + f(y)$. Form a tree S' by removing x and y from T' , and replacing their parent by z' .
 - S' is a tree code on $n - 1$ letters.
- $cost(T') = cost(S') + f(x) + f(y) \geq cost(S) + f(x) + f(y) = cost(T)$.
 - First equality because x, y at depth one greater than z .
 - First inequality because S is opt tree code on $n - 1$ letters.
- So, the tree T produced by Huffman encoding is optimal.