



Amortized analysis, Fibonacci heaps

CS240

Spring 2023

Rui Fan



Amortized analysis

- Suppose we want to bound the amount of time to perform n (possibly different) operations on a data structure.
- If max amount of time for each operation is $f(n)$, $n \cdot f(n)$ is upper bound on the time for all the operations.
- But some operations might take more time than others.
 - Even the same operation can take different amounts of time each time it's executed.
 - So $n \cdot f(n)$ may overestimate actual amount of time taken.
 - The bound isn't tight.
- Amortized analysis looks at the average amount of time for each operation over all the operations.
 - The average is taken over the worst case execution, i.e. a sequence of operations with the highest average cost for the data structure.

Potential method

- To keep track of the true total cost of a sequence of operations, we use a potential function $\Phi: D \rightarrow \mathbb{R}$, where D is the set of states of the data structure.
- Let D_i be the state of the data structure after applying the i 'th operation, and c_i be the cost of the i 'th operation.
- **Def** The amortized cost for the i 'th operation is $\hat{c}_i = c_i + \Phi(D_i) - \Phi(D_{i-1})$.
- Using the amortized cost, we sometimes overcharge and sometimes undercharge for operations.
 - I.e. when $\hat{c}_i > c_i$, we overcharge, and when $\hat{c}_i < c_i$ we undercharge.
- However, the total amortized cost is at least the total actual cost, i.e. $\sum_i \hat{c}_i \geq \sum_i c_i$.
 - So total amortized cost is an upper bound on total actual cost.
- If we design the right potential function, we can keep track of the total cost by tracking the amortized costs.
 - The amortized cost is sometimes easier to analyze than directly keeping track of actual costs.
 - This leads to tight bounds for many data structures.

Potential method

- When we overcharge, i.e. $\hat{c}_i = c_i + \Phi(D_i) - \Phi(D_{i-1}) > c_i$, Φ increases.
 - We “store” the extra amortized cost $\hat{c}_i - c_i$ we charged the i 'th operation in Φ
 - Φ is also called “credit” or “potential” (energy).
- When we undercharge, i.e. $\hat{c}_i < c_i$, Φ decreases.
 - We use some of the stored credit to pay for the $c_i - \hat{c}_i$ amount of actual cost that the amortized cost doesn't account for.
- **Lemma** Suppose $\Phi(D_n) \geq \Phi(D_0)$. Then $\sum_{i=1}^n \hat{c}_i \geq \sum_{i=1}^n c_i$.
- **Proof**
$$\sum_{i=1}^n \hat{c}_i = \sum_{i=1}^n (c_i + \Phi(D_i) - \Phi(D_{i-1})) = (\sum_{i=1}^n c_i) + \Phi(D_n) - \Phi(D_0) \geq \sum_{i=1}^n c_i.$$
 - The second equality follows because all the terms except $\Phi(D_n), \Phi(D_0)$ telescope away.
- A simple way to ensure $\Phi(D_n) \geq \Phi(D_0)$ is to design Φ so that $\Phi(D_0) = 0$, and $\Phi(D_i) \geq 0$ for all i .

Example: Binary counter

- Consider a k-digit binary counter. When we increment the counter, we flip some bits.
 - Suppose each bit flip costs 1 unit.
- What is the total cost for incrementing the counter n times, starting from 0?
- Since there are k digits, a trivial upper bound is $O(nk)$.
- However, the actual number of bit flips is much less, because most increments only flip a few bits.
- We use the potential method to show the total cost is $O(n)$.
 - In fact, it's at most $2n$.

Counter value	A[7]	A[6]	A[5]	A[4]	A[3]	A[2]	A[1]	A[0]	Total cost
0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	1	1
2	0	0	0	0	0	0	1	0	3
3	0	0	0	0	0	0	1	1	4
4	0	0	0	0	0	1	0	0	7
5	0	0	0	0	0	1	0	1	8
6	0	0	0	0	0	1	1	0	10
7	0	0	0	0	0	1	1	1	11
8	0	0	0	0	1	0	0	0	15
9	0	0	0	0	1	0	0	1	16
10	0	0	0	0	1	0	1	0	18
11	0	0	0	0	1	0	1	1	19
12	0	0	0	0	1	1	0	0	22
13	0	0	0	0	1	1	0	1	23
14	0	0	0	0	1	1	1	0	25
15	0	0	0	0	1	1	1	1	26
16	0	0	0	1	0	0	0	0	31



Example: Binary counter

- Let $\Phi(D_i) = b_i$, where D_i is the state of the counter after i increments, and b_i is the number of 1's in D_i .
- Suppose the i 'th operation sets t_i bits from 1 to 0.
 - Then the actual cost is $c_i = t_i + 1$.
 - This sets t_i bits from 1 to 0, and one bit from 0 to 1 for the carry.
- If $b_i = 0$, then the i 'th operation reset all the bits.
 - Also, all the bits were set in D_{i-1} .
 - So $t_i = b_{i-1} = k$.
- If $b_i > 0$, then $b_i = b_{i-1} - t_i + 1$.
 - t_i bits went from 1 to 0, and one carry bit went from 0 to 1.
- In both cases, $b_i \leq b_{i-1} - t_i + 1$.



Example: Binary counter

- Since $b_i \leq b_{i-1} - t_i + 1$, then $\Phi(D_i) - \Phi(D_{i-1}) = b_i - b_{i-1} \leq 1 - t_i$.
- So the amortized cost $\hat{c}_i = c_i + \Phi(D_i) - \Phi(D_{i-1}) \leq (t_i + 1) + (1 - t_i) = 2$.
- Finally, $\Phi(D_0) = 0$, since the counter is initially 0, and $\Phi(D_n) \geq 0$.
- Thus, by the lemma the total cost for all n increments is $\sum_{i=1}^n c_i \leq \sum_{i=1}^n \hat{c}_i \leq 2n$.

Fibonacci heaps

- A Fibonacci heap is a type of heap that implements certain operations faster than a binary heap, in amortized time.
 - The time complexities of the circled operations are amortized. The rest are worst case.
- It can be used to speed up a number of graph algorithms asymptotically.
- Both Dijkstra's and Prim's algorithms take $O((V+E) \log V)$ time with a binary heap, and $O(E + V \log V)$ time with a Fibonacci heap.
 - Both algorithms decrease the key value heap items $O(E)$ times.
 - This takes $O(E \log V)$ time on a binary heap, and $O(E)$ time on a Fibonacci heap.
- Fibonacci heaps are more complicated than binary heaps, and often don't perform better in practice.
- When decreasing or deleting a key, assume we have a pointer to the node with the key.
 - Otherwise finding the node takes $O(n)$ time, where n is the number of items.

Procedure	Binary heap (worst-case)	Fibonacci heap (amortized)
MAKE-HEAP	$\Theta(1)$	$\Theta(1)$
INSERT	$\Theta(\lg n)$	$\Theta(1)$
MINIMUM	$\Theta(1)$	$\Theta(1)$
EXTRACT-MIN	$\Theta(\lg n)$	$O(\lg n)$
UNION	$\Theta(n)$	$\Theta(1)$
DECREASE-KEY	$\Theta(\lg n)$	$\Theta(1)$
DELETE	$\Theta(\lg n)$	$O(\lg n)$

DIJKSTRA(G, w, s)

```

1 INITIALIZE-SINGLE-SOURCE( $G, s$ )
2  $S = \emptyset$ 
3  $Q = G.V$ 
4 while  $Q \neq \emptyset$ 
5    $u = \text{EXTRACT-MIN}(Q)$ 
6    $S = S \cup \{u\}$ 
7   for each vertex  $v \in G.Adj[u]$ 
8     RELAX( $u, v, w$ )

```

RELAX(u, v, w)

```

1 if  $v.d > u.d + w(u, v)$ 
2    $v.d = u.d + w(u, v)$ 
3    $v.\pi = u$ 

```

MST-PRIM(G, w, r)


```

1 for each  $u \in G.V$ 
2    $u.key = \infty$ 
3    $u.\pi = \text{NIL}$ 
4  $r.key = 0$ 
5  $Q = G.V$ 
6 while  $Q \neq \emptyset$ 
7    $u = \text{EXTRACT-MIN}(Q)$ 
8   for each  $v \in G.Adj[u]$ 
9     if  $v \in Q$  and  $w(u, v) < v.key$ 
10       $v.\pi = u$ 
11       $v.key = w(u, v)$ 

```


Structure of Fibonacci heap


- A Fibonacci heap H consists of a set of rooted trees.
- Each tree satisfies the min heap property, i.e. each node's key is less than those of all its children.
- The trees are linked in a doubly-linked root list.
 - These roots are connected by the dashed line in the top figure.
- The minimum node is a root, and is pointed to by $H.min$.
- $H.n$ stores the total number of nodes in all trees.
- Within each tree, the nodes at each level are also linked in a doubly-linked list.
- The two figures show the same Fibonacci heap, but the top figure avoids showing the linked list for clarity.



Potential function



- Some nodes are marked.
 - These are shown in black.
- Marks are only used during decrease key and deletion operations, and are described later.
 - They help ensure each node has a large number of children.
 - This ensures each tree is not too tall, and so each operation is fast.
- Suppose H has $t(H)$ root nodes and $m(H)$ marked nodes.
- The potential of H is

$$\Phi(H) = t(H) + 2m(H)$$



Basic operations

- Let H and H' denote the heap before and after an operation.
- **Make-Heap**
 - Make an empty heap. Set $H'.n=0$, $H'.min=NIL$.
 - Cost = $O(1)$.
 - Amortized cost = $O(1)$, since $\Phi(H) = \Phi(H') = 0$.
- **Insert a node**
 - Add the new node to the root list, left of the min node.
 - Change $H.min$ if new node's key is smaller.
 - Cost = $O(1)$.
 - Amortized cost = $O(1)$.
 - Number of roots increases by 1.
 - So $\Phi(H') - \Phi(H) = (t(H) + 1 + 2m(H)) - (t(H) + 2m(H)) = O(1)$.



FIB-HEAP-INSERT(H, x)

```

1   $x.degree = 0$ 
2   $x.p = NIL$ 
3   $x.child = NIL$ 
4   $x.mark = FALSE$ 
5  if  $H.min == NIL$ 
6      create a root list for  $H$  containing just  $x$ 
7       $H.min = x$ 
8  else insert  $x$  into  $H$ 's root list
9      if  $x.key < H.min.key$ 
10          $H.min = x$ 
11   $H.n = H.n + 1$ 
  
```

Basic operations

■ Find the minimum

- Return $H.min$.
- Cost = amortized cost = $O(1)$.

■ Union of two heaps

- Concatenate the root lists of the two heaps H_1 and H_2 .
- Set $H.min$ to $\min(H_1.min, H_2.min)$
- Cost = $O(1)$.
- Since new root list is the union of the two old root lists, the change in potential is $\Phi(H') - (\Phi(H_1) + \Phi(H_2)) = t(H') + 2m(H') - (t(H_1) + 2m(H_1) + t(H_2) + 2m(H_2)) = 0$.
- Thus the amortized cost = $O(1)$.

FIB-HEAP-UNION(H_1, H_2)

```
1   $H = \text{MAKE-FIB-HEAP}()$ 
2   $H.min = H_1.min$ 
3  concatenate the root list of  $H_2$  with the root list of  $H$ 
4  if ( $H_1.min == \text{NIL}$ ) or ( $H_2.min \neq \text{NIL}$  and  $H_2.min.key < H_1.min.key$ )
5       $H.min = H_2.min$ 
6   $H.n = H_1.n + H_2.n$ 
7  return  $H$ 
```



Extract min node

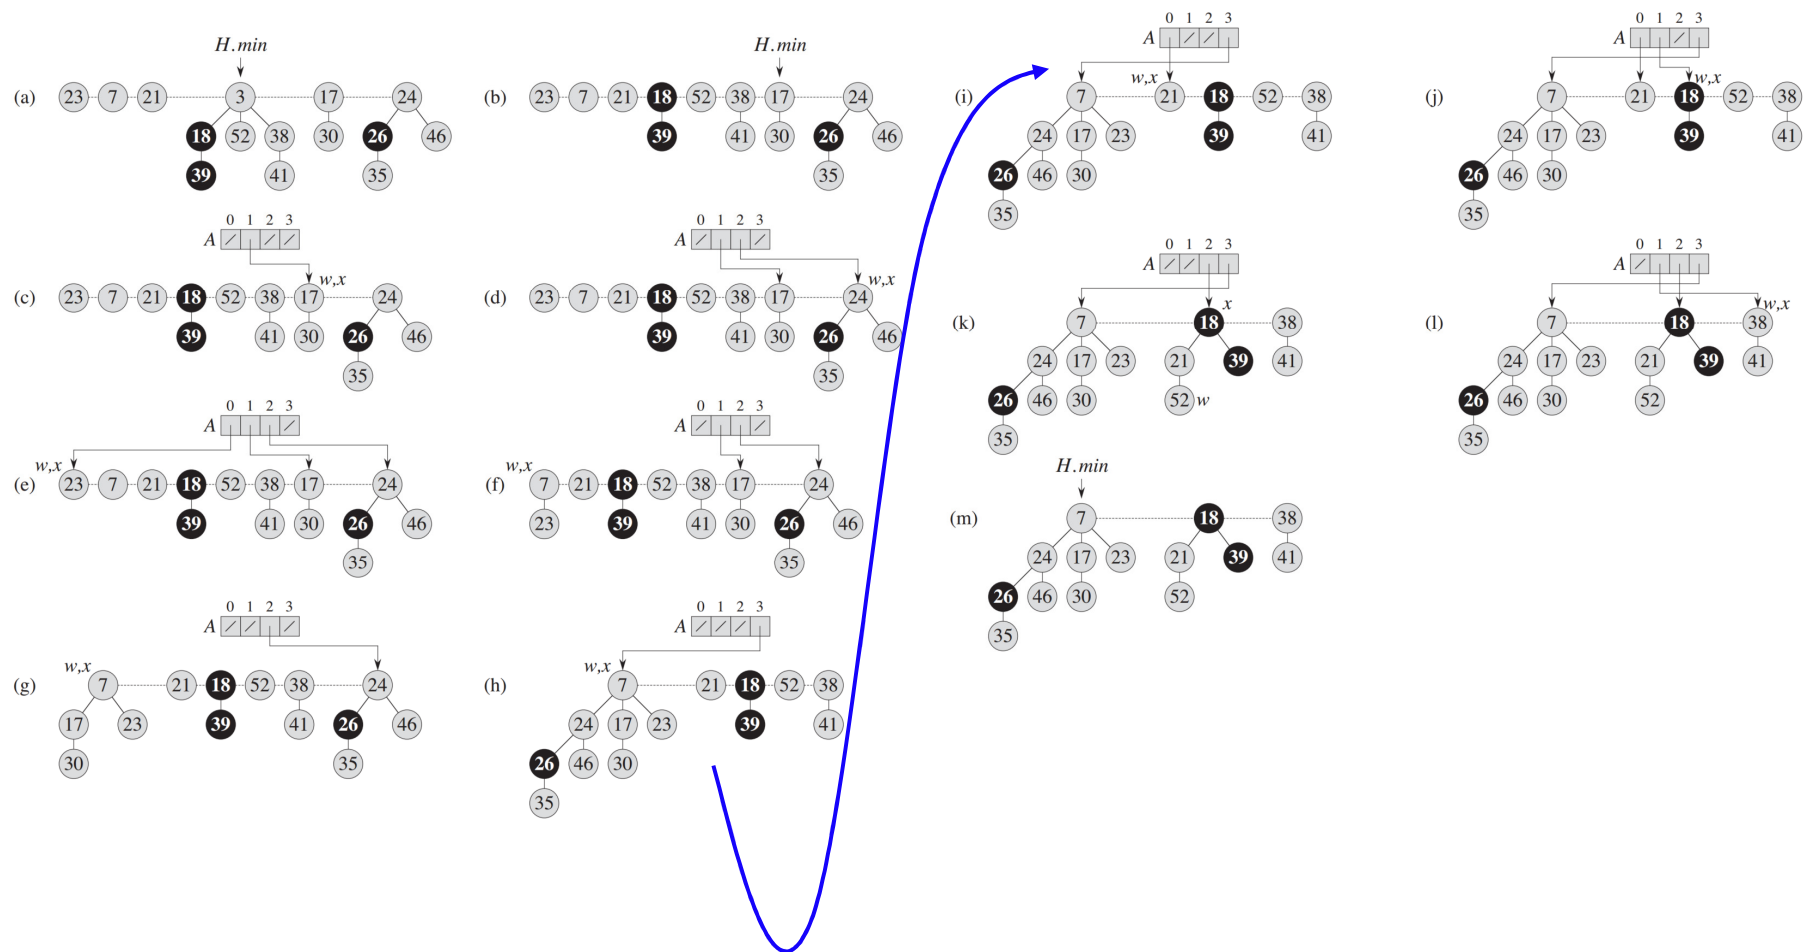
- First remove the H.min node.
- Then add each of its children (along with its subtree) to the root list.
 - There may now be many trees in the root list.
 - To find the new H.min, we need to iterate through the roots of all the trees, which may be slow.
 - So we want to decrease the number of trees in the root list.
- **Def** The degree of a node is its number of children.
- We merge some of the trees in the root list, so that none of the roots have the same degree.
 - The merging function is called CONSOLIDATE.



CONSOLIDATE

- Let $D(H.n)$ be upper bound on the degree of any node in a Fibonacci heap with n nodes.
 - We show later $D(H.n) = O(\log n)$.
- Use an array A of size $D(H.n)+1$.
 - $A[i]$ points to a tree in the root list with degree i .
- Iterate through all the trees in the root list.
 - If the current tree x we process has degree d , and $A[d] = y \neq NIL$, then there's already a tree y in the root list with degree d .
 - Since we don't want two trees in the root list with the same degree, we link the roots of x and y .
 - This creates a tree in the root list with degree $d+1$, and removes the tree with degree d .
 - The direction we link depends on which root has the smaller key.
 - Then set $A[d]=NIL$, and set $A[d+1]$ to point to newly linked tree.
 - If the new root is marked, clear the mark.
- Finally, iterate through A array, and set $H.min$ to the min root value.

Example: Extract min



Pseudocode for extract min

FIB-HEAP-EXTRACT-MIN(H)

```
1   $z = H.min$ 
2  if  $z \neq \text{NIL}$ 
3      for each child  $x$  of  $z$ 
4          add  $x$  to the root list of  $H$ 
5           $x.p = \text{NIL}$ 
6  remove  $z$  from the root list of  $H$ 
7  if  $z == z.right$ 
8       $H.min = \text{NIL}$ 
9  else  $H.min = z.right$ 
10  CONSOLIDATE( $H$ )
11   $H.n = H.n - 1$ 
12  return  $z$ 
```

FIB-HEAP-LINK(H, y, x)

```
1  remove  $y$  from the root list of  $H$ 
2  make  $y$  a child of  $x$ , incrementing  $x.degree$ 
3   $y.mark = \text{FALSE}$ 
```

CONSOLIDATE(H)

```
1  let  $A[0 \dots D(H.n)]$  be a new array
2  for  $i = 0$  to  $D(H.n)$ 
3       $A[i] = \text{NIL}$ 
4  for each node  $w$  in the root list of  $H$ 
5       $x = w$ 
6       $d = x.degree$ 
7      while  $A[d] \neq \text{NIL}$ 
8           $y = A[d]$  // another node with the same degree as  $x$ 
9          if  $x.key > y.key$ 
10              exchange  $x$  with  $y$ 
11          FIB-HEAP-LINK( $H, y, x$ )
12           $A[d] = \text{NIL}$ 
13           $d = d + 1$ 
14       $A[d] = x$ 
15   $H.min = \text{NIL}$ 
16  for  $i = 0$  to  $D(H.n)$ 
17      if  $A[i] \neq \text{NIL}$ 
18          if  $H.min == \text{NIL}$ 
19              create a root list for  $H$  containing just  $A[i]$ 
20               $H.min = A[i]$ 
21          else insert  $A[i]$  into  $H$ 's root list
22              if  $A[i].key < H.min.key$ 
23                   $H.min = A[i]$ 
```


Complexity for extract min


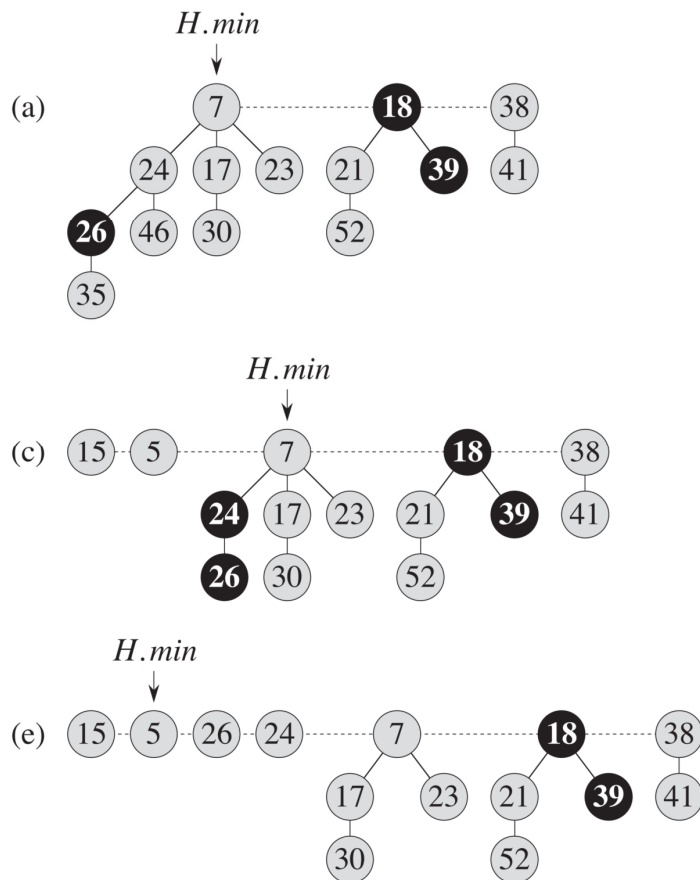
- Let H denote the heap before the extract min.
- The real cost includes
 - $O(D(n))$ for moving children of $H.min$ to root list.
 - The for loop in lines 4-14 of CONSOLIDATE operate on a list of size at most $D(n)+t(H)-1$.
 - Every time through the while loop in lines 7-13, we link two of the trees in the root list.
 - Each tree can be linked (to a tree whose root has a smaller key) at most once.
 - So the total number of iterations of the while loop is at most the root list size, i.e. $O(D(n)+t(H))$.
 - So the real cost is $O(D(n)+t(H))$.
- For the amortized cost, the potential before the extract min is at most $t(H)+2m(H)$.
- The potential after extract is $\leq (D(n)+1)+2m(H)$.
 - All trees in root list of H' have different degrees, and max degree is $D(n)$.
 - No new nodes get marked during extract.
- So amortized cost is
$$O(D(n) + t(H)) + ((D(n) + 1) + 2m(H)) - (t(H) + 2m(H))$$
$$= O(D(n)) + O(t(H)) - t(H) = O(D(n)).$$
 - The last equality follows because we can scale up the units of the potential to cancel out the hidden constant in $O(t(H))$.



Decreasing key and marking

- To decrease a node x 's key, check if the new key violates the heap property.
 - If not, we're done.
 - Otherwise, move x and its subtree to the root list.
 - We say we cut out x (and its subtree).
 - Unmark x , if it's marked.
- A node is marked if one of its children has been cut, since the last time it's been cut.
- The second time a node's children is cut out, we move the node (and its subtree) to the root list and unmark it.
- Let y be x 's parent.
 - If y is not marked, mark y , since we cut one of its children.
 - If y is already marked, move y and its subtree to the root list, and then unmark y .
 - Let z be y 's parent.
 - If z is not marked, stop. Otherwise, cut z and move it to the root list, and repeat the previous steps for z 's parent, etc.
- One decrease key can create a sequence of cascading cuts.

Example: decrease key



- (a) shows the original Fibonacci heap.
- (b) shows the heap after node 46 is decreased to 15.
- (c)-(e) show the cascading cuts after node 35 is decreased to 5.

Pseudocode for decrease key and delete

FIB-HEAP-DECREASE-KEY(H, x, k)

```
1  if  $k > x.key$ 
2      error "new key is greater than current key"
3   $x.key = k$ 
4   $y = x.p$ 
5  if  $y \neq \text{NIL}$  and  $x.key < y.key$ 
6      CUT( $H, x, y$ )
7      CASCADING-CUT( $H, y$ )
8  if  $x.key < H.min.key$ 
9       $H.min = x$ 
```

FIB-HEAP-DELETE(H, x)

```
1  FIB-HEAP-DECREASE-KEY( $H, x, -\infty$ )
2  FIB-HEAP-EXTRACT-MIN( $H$ )
```

CUT(H, x, y)

```
1  remove  $x$  from the child list of  $y$ , decrementing  $y.degree$ 
2  add  $x$  to the root list of  $H$ 
3   $x.p = \text{NIL}$ 
4   $x.mark = \text{FALSE}$ 
```

CASCADING-CUT(H, y)

```
1   $z = y.p$ 
2  if  $z \neq \text{NIL}$ 
3      if  $y.mark == \text{FALSE}$ 
4           $y.mark = \text{TRUE}$ 
5      else CUT( $H, y, z$ )
6      CASCADING-CUT( $H, z$ )
```

To delete a key, simply decrease its value to $-\infty$ and then do a extract-min.

Complexity for decrease key

- Let H denote the heap before the decrease key operation.
- Cutting out a node takes $O(1)$ time.
- Suppose a decrease key operation creates c cascading cuts.
- Then the actual cost is $O(c)$.
- For the amortized cost
 - Each cut creates one more tree in the root list.
 - It also removes one marked node.
 - After the decrease key, the root list contains $t(H)+c$ trees.
 - It also contains $\leq m(H) - c + 2$ marked nodes.
 - $c-1$ nodes were unmarked by cascading cuts, and the last call to CASCADING-CUT may have marked a node.
 - So the change in potential is $\left((t(H)) + c + 2(m(H) - c + 2) \right) - ((t(H) + 2m(H))) = 4 - c$.
 - So the amortized cost is $O(c) + 4 - c = O(1)$ by scaling the hidden constant in the potential appropriately.

Bounding the max degree

- So far, all the operations have $O(1)$ amortized cost, except extract-min (and delete, which calls extract-min).
- extract-min has amortized cost $O(D(n))$, where $D(n)$ is the max degree of any node in the Fibonacci heap with n nodes.
- **Def** The golden ratio is $\phi = \frac{1+\sqrt{5}}{2} \approx 1.618$.
 - ϕ is the positive solution to the equation $x^2 = x + 1$.
- Recall the Fibonacci F_n is defined by $F_0 = 0, F_1 = 1$, and $F_n = F_{n-1} + F_{n-2}$ for $n \geq 2$.
 - The sequence is 0, 1, 1, 2, 3, 5, 8, 13, 21, 34, ...
- **Fact 1** $F_n = \lfloor \frac{\phi^n}{\sqrt{5}} + \frac{1}{2} \rfloor$.
 - For a proof, see section 3.2 of *Introduction to Algorithms*.
- **Fact 2** $F_{n+2} = 1 + \sum_{i=0}^n F_i$.
- **Fact 3** $F_{n+2} \geq \phi^n$.
- We show $D(n) \leq \lfloor \log_\phi n \rfloor$.
- **Def** For any node let $x.\text{deg}$ denote its degree, and $\text{size}(x)$ be the number of nodes in x 's subtree (including x).

Bounding the max degree

- **Lemma 1** Let x be a node in a Fibonacci heap, and suppose $x.deg = k$. Let y_1, \dots, y_k be the children of x , in the order they were linked to x , from earliest to latest. Then $y_1.deg \geq 0$, and $y_i.deg \geq i - 2$ for $i = 2, \dots, k$.
- **Proof** Obviously $y_1.deg \geq 0$.
 - For $i \geq 2$, when y_i was linked to x , y_1, \dots, y_{i-1} were already children of x , and so x had degree $\geq i - 1$.
 - y_i was linked to x during CONSOLIDATE.
 - So when y_i was linked, we had $y_i.deg = x.deg \geq i - 1$.
 - Since y_i was linked to x , it could have lost at most one child.
 - As soon as y_i loses two children, it's cut and moved to the root list.
 - So $y_i.deg \geq i - 2$.

Bounding the max degree

- **Lemma 2** Let x be a node in a Fibonacci heap, and suppose $x.deg = k$. Then $size(x) \geq F_{k+2} \geq \phi^k$.
- **Proof** We use induction on k . The bound holds for $k = 0, 1$. For higher k , let y_1, \dots, y_k denote the children of x .
 - By Lemma 1, $y_i.deg \geq i - 2$ for $i \geq 2$.
 - So by induction, $size(y_i) \geq F_i$, for $i \geq 2$.
 - Also, $size(y_0), size(y_1) \geq 1$.
 - We have $size(x) \geq \sum_{i=0}^k size(y_i) \geq 2 + \sum_{i=2}^k size(y_i) \geq 2 + \sum_{i=2}^k F_i = 1 + \sum_{i=0}^k F_i = F_{k+2} \geq \phi^k$.
 - The last two equalities follow by Facts 2 and 3.
- **Cor** For any n node Fibonacci heap H , the max degree $D(n) = O(\log n)$.
- **Proof** Let x be any node in H , and let $k = x.deg$.
 - We have $n \geq size(x) \geq \phi^k$.
 - So $k \leq \lfloor \log_{\phi} n \rfloor$.