

Announcement

- ▶ Homework 2
 - ▶ Available in Blackboard -> Homework
 - ▶ Due: Mar. 9, 11:59pm (one week)



Sequence to Sequence

SLP3 Ch 10; INLP Ch 18

Sequence to sequence

- ▶ Input a sequence, output another sequence
 - ▶ Often abbreviated to seq2seq
- ▶ Many applications
 - ▶ Machine translation (the most well-known seq2seq task)
 - ▶ Input: source language text
 - ▶ Output: target language text
 - ▶ Paraphrase
 - ▶ Input: a sentence
 - ▶ Output: a restatement of the meaning using other words
 - ▶ Style transfer
 - ▶ Input: some text
 - ▶ Output: a restatement in a specified style (e.g., formal writing, advertisement, positive sentiment, Shakespearean writing style, etc.)



Sequence to sequence

- ▶ Input a sequence, output another sequence
 - ▶ Often abbreviated to seq2seq
- ▶ Many applications
 - ▶ Summarization
 - ▶ Input: a document
 - ▶ Output: a few sentences that summarize the document
 - ▶ Question answering
 - ▶ Input: a question
 - ▶ Output: the answer
 - ▶ Dialog
 - ▶ Input: previous utterances
 - ▶ Output: next utterance




Sequence to sequence

► More applications

- **Input:** Tolkien's epic novel The Lord of the Rings was published in 1954-1955, years after the book was completed.
- **Joint entity and relation extraction**
 - [Tolkien | *person*]'s epic novel [The Lord of the Rings | *book* | *author* = Tolkien] was published in 1954-1955, years after the book was completed.
- **Semantic role labeling**
 - Tolkien's epic novel [The Lord of the Rings | *subject*] [was published | *predicate*] [in 1954-1955 | *temporal*], years after the book was completed.
- **Coreference resolution**
 - [Tolkien]'s epic novel [The Lord of the Rings] was published in 1954-1955, years after the [book | The Lord of the Rings] was completed.
- More on these tasks later...

Comparing Methods of Machine Translation



Sources: http://www.meta-net.eu/events/meta-forum-2016/slides/09_sennrich.pdf & <http://matrix.statmt.org/>

SMT: statistical machine translation

NMT: neural machine translation \Leftarrow neural seq2seq



NMT: a big success story of NLP + deep learning

- ▶ Neural Machine Translation went from **a fringe research attempt** in 2014 to the **leading standard method** in 2016
 - ▶ 2014: First seq2seq paper published
 - ▶ 2016: Google Translate switches from SMT to NMT - and by 2018 everyone has



- ▶ SMT systems, built by **hundreds of engineers** over **many years**, outperformed by NMT systems trained by a small group of engineers in a few months





Neural Seq2Seq Methods

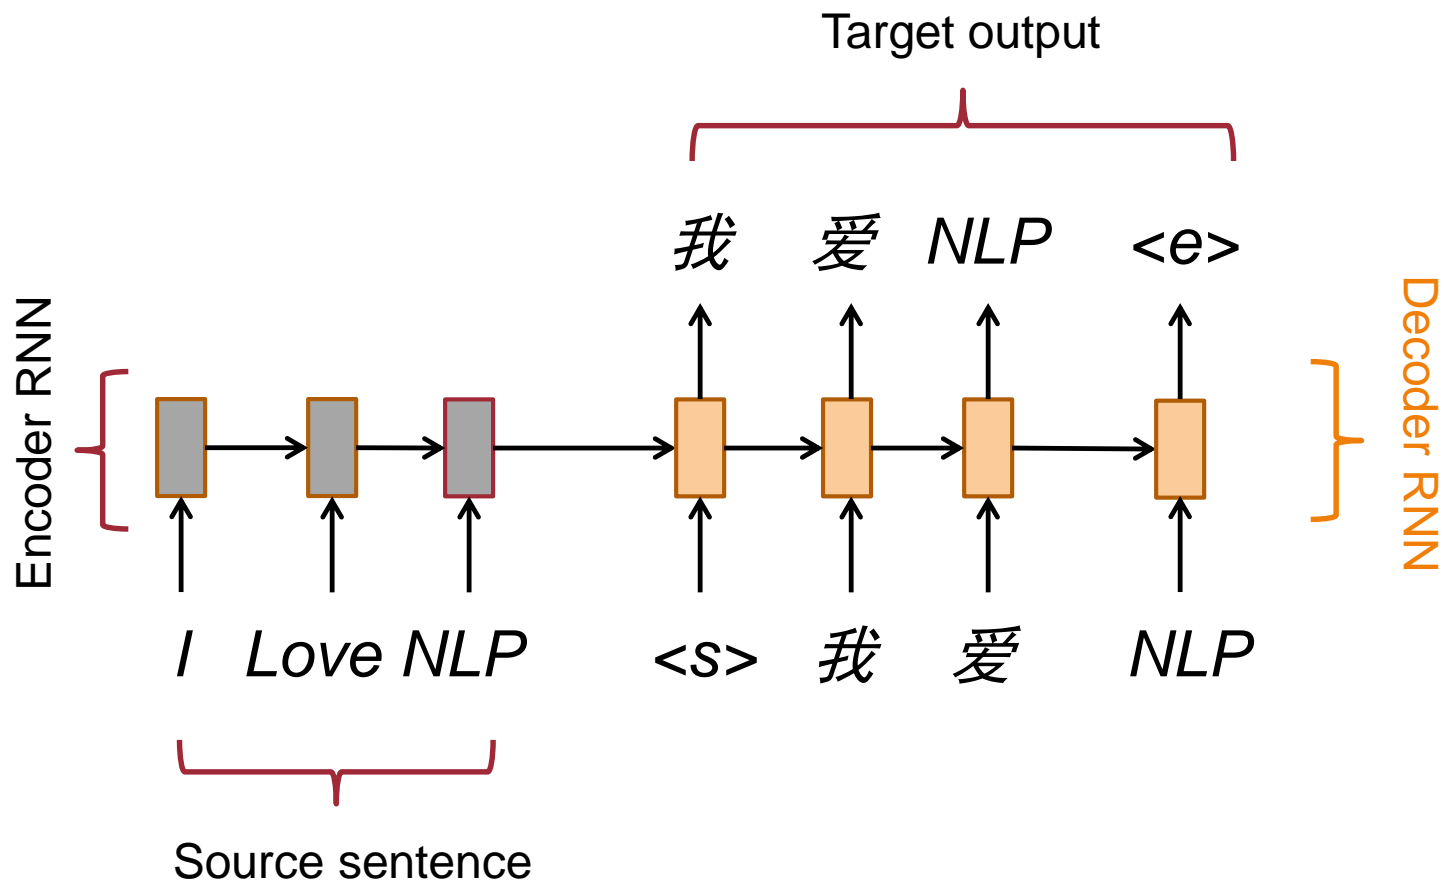


Neural Seq2seq Methods

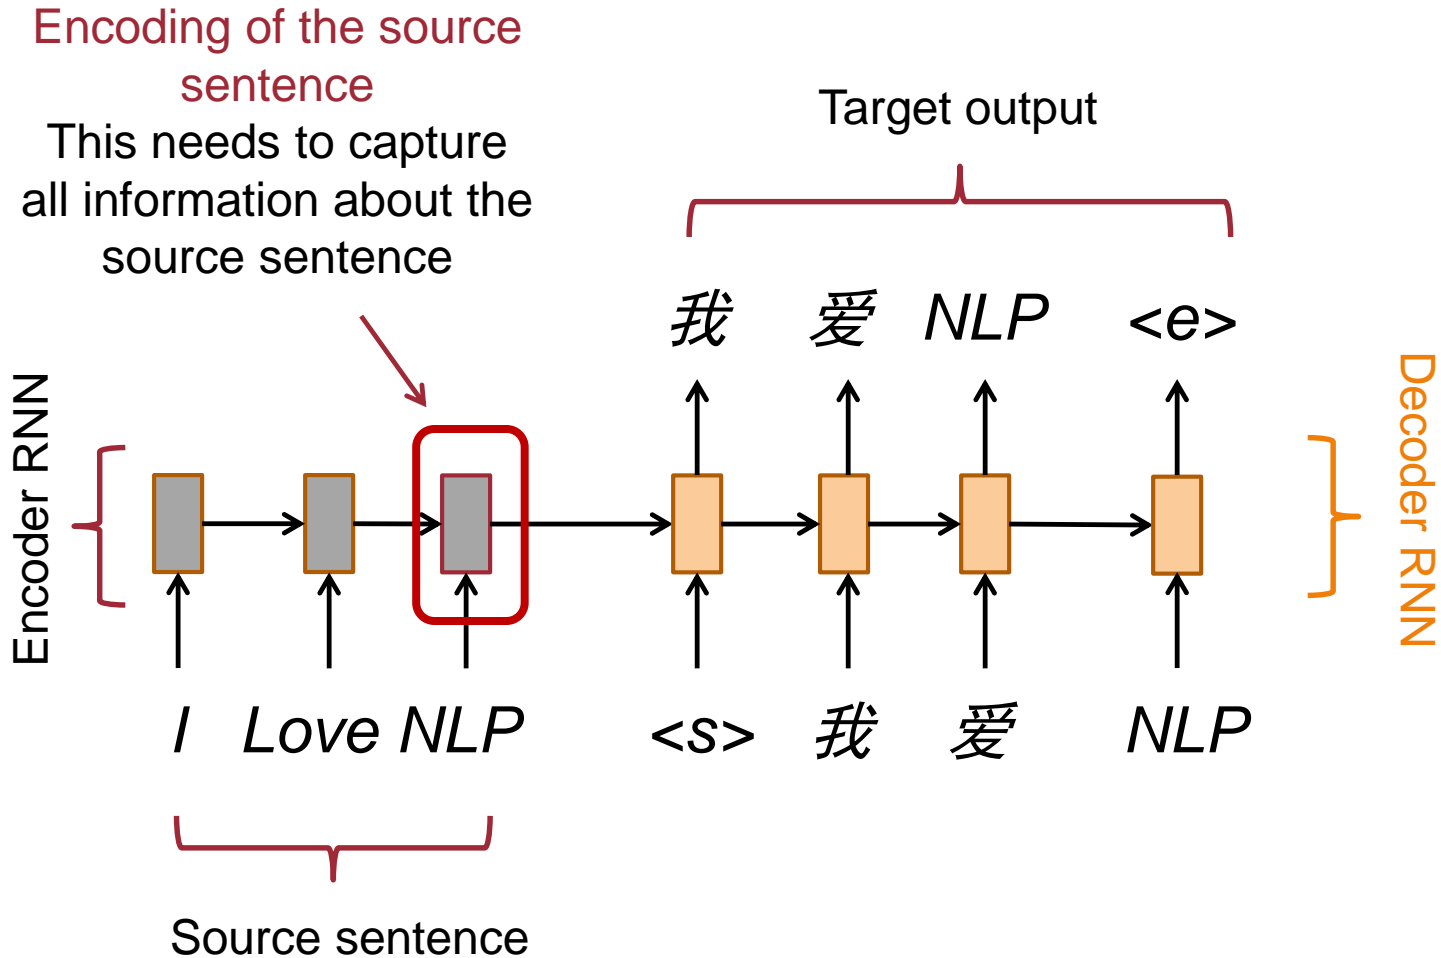
- ▶ Similar to neural language modeling methods
 - ▶ Recurrent neural network
 - ▶ Attention
 - ▶ Transformer
- ▶ Difference: now we have an encoder to process the input sequence and a decoder to produce the output sequence



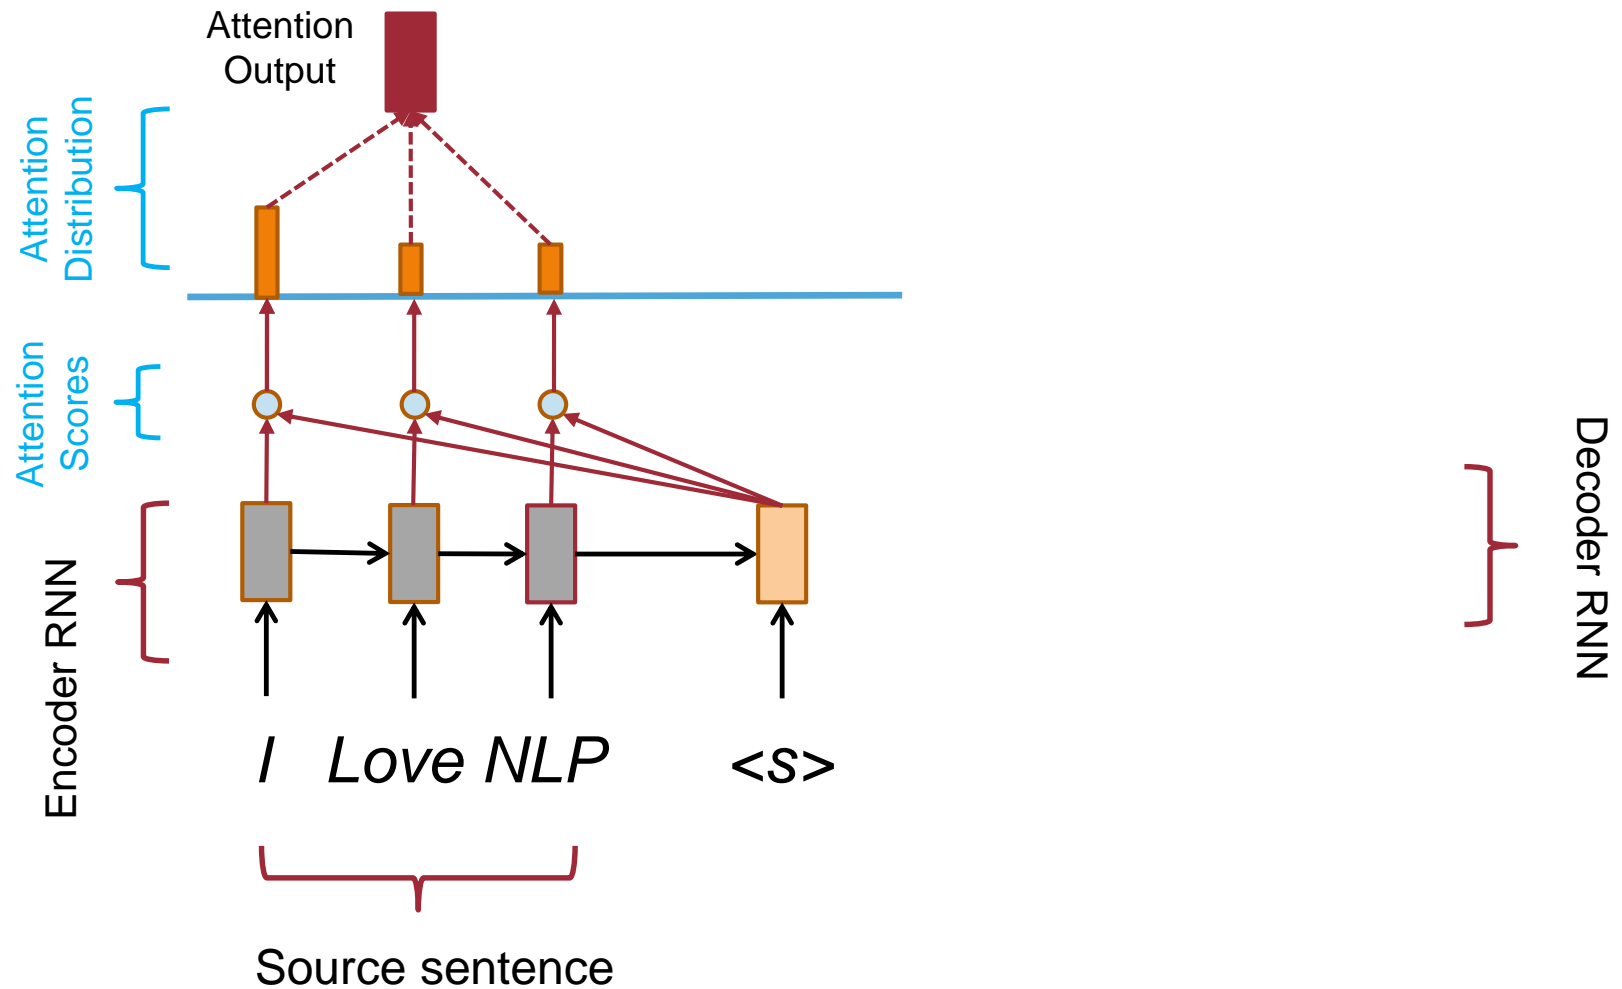
RNN



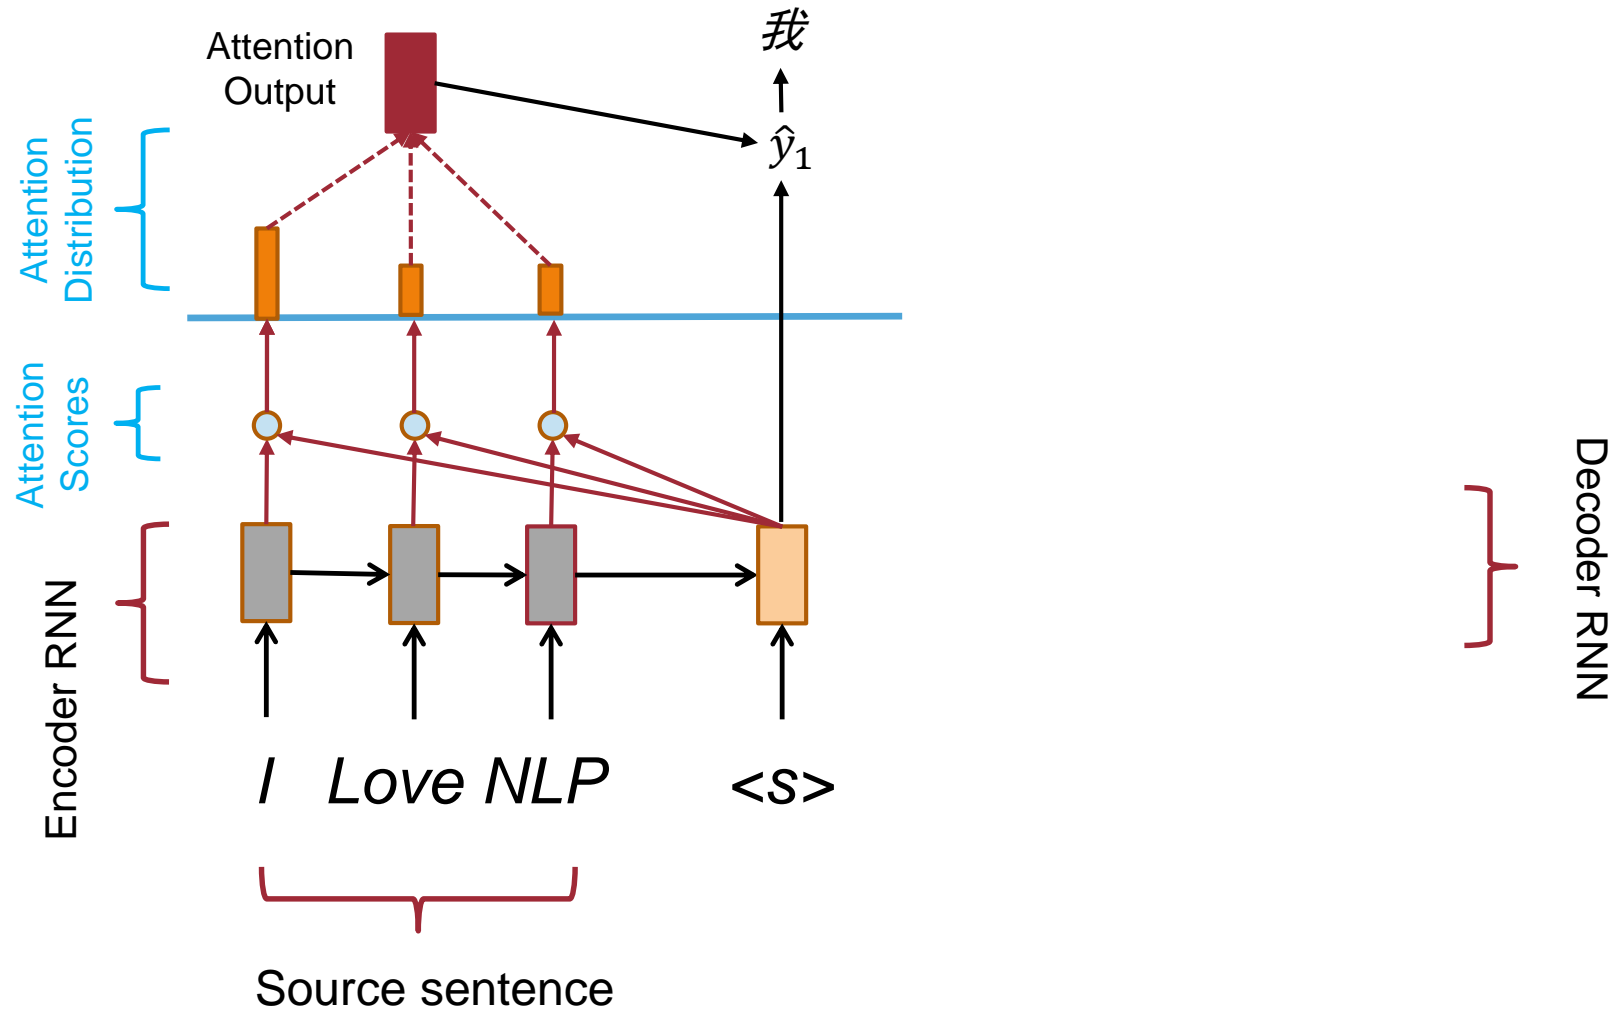
RNN: the bottleneck problem



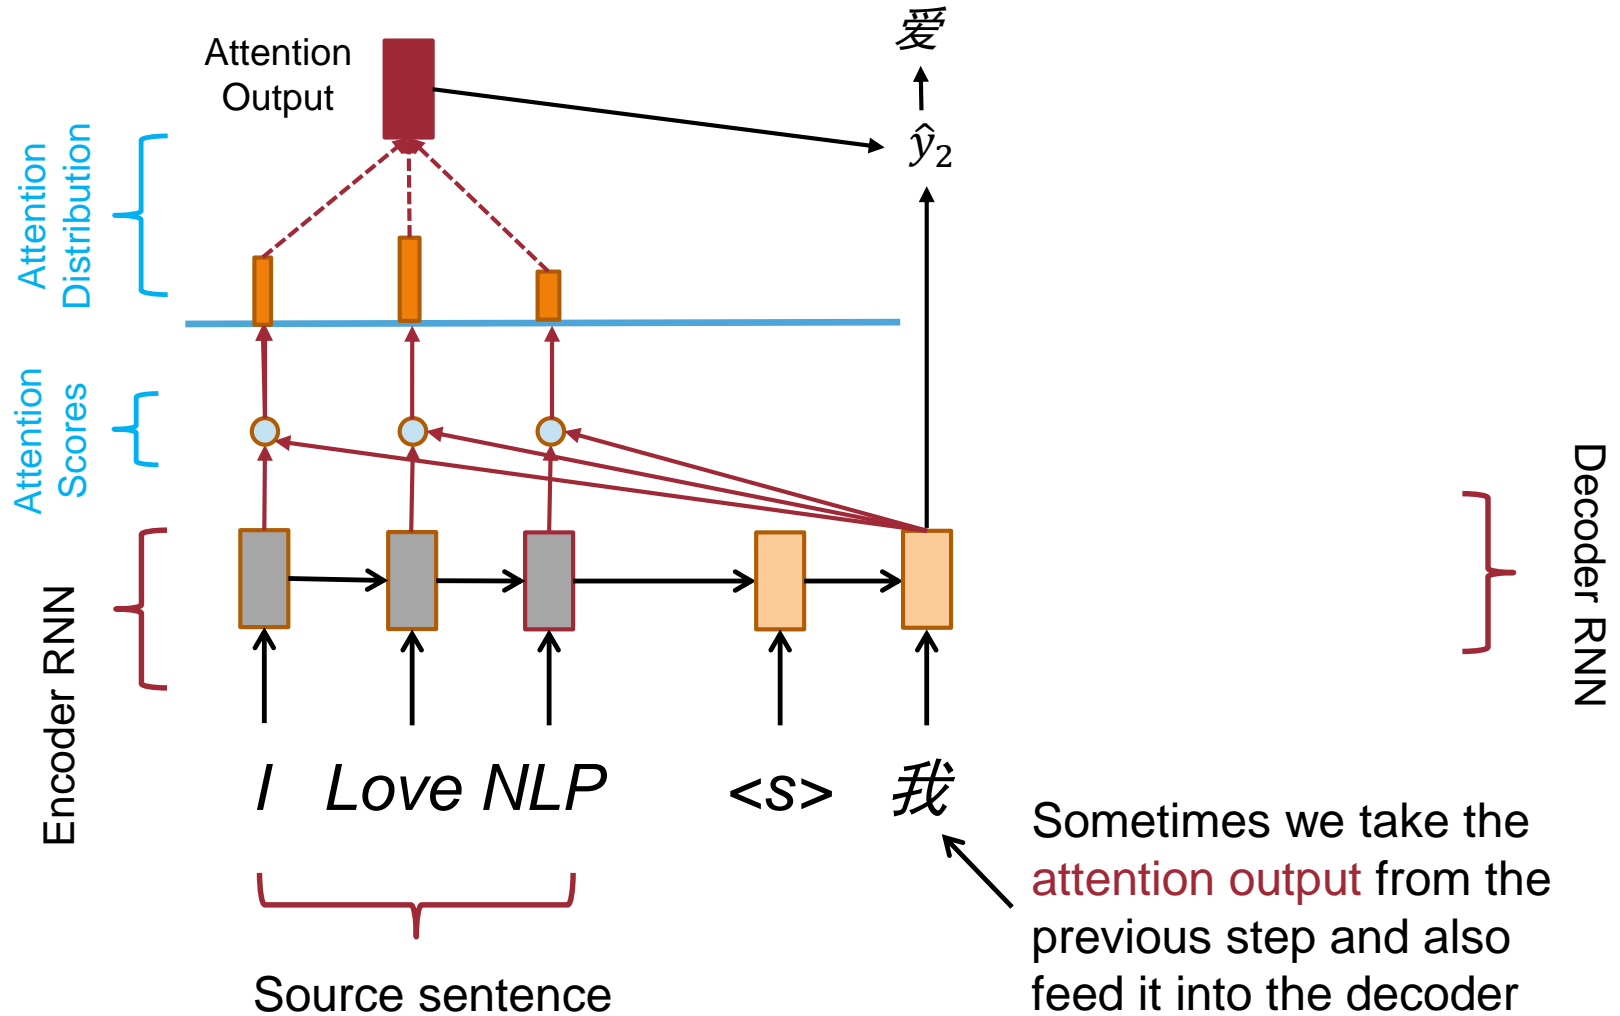
Seq2seq with attention



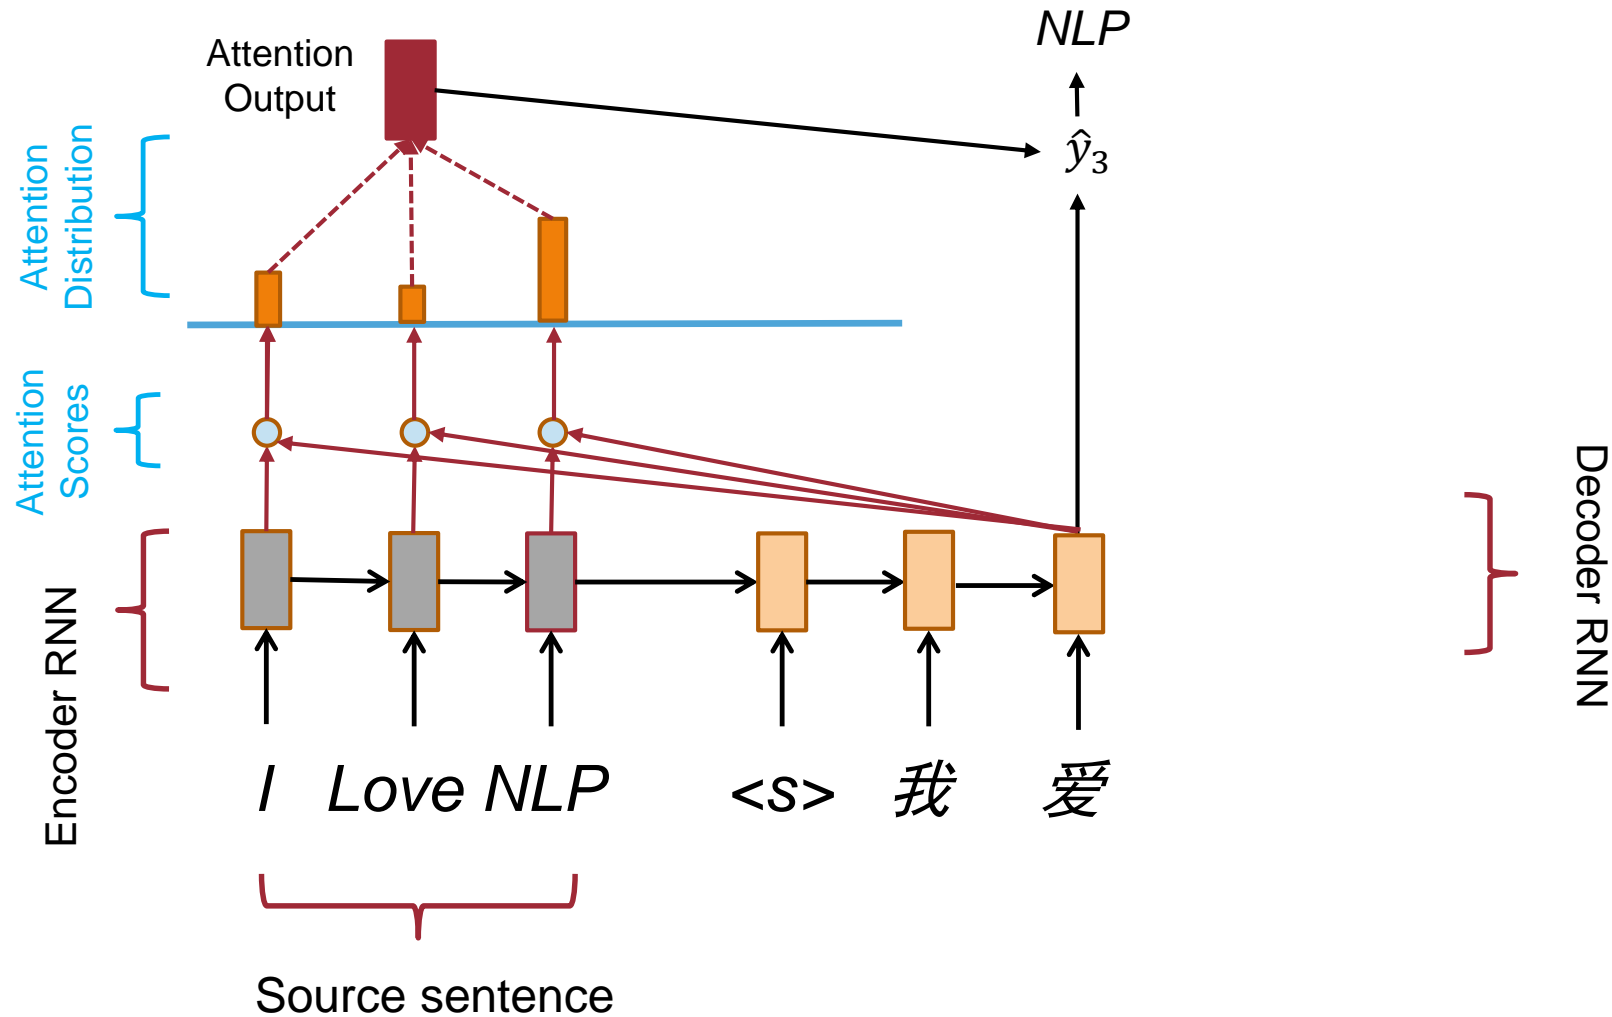
Seq2seq with attention



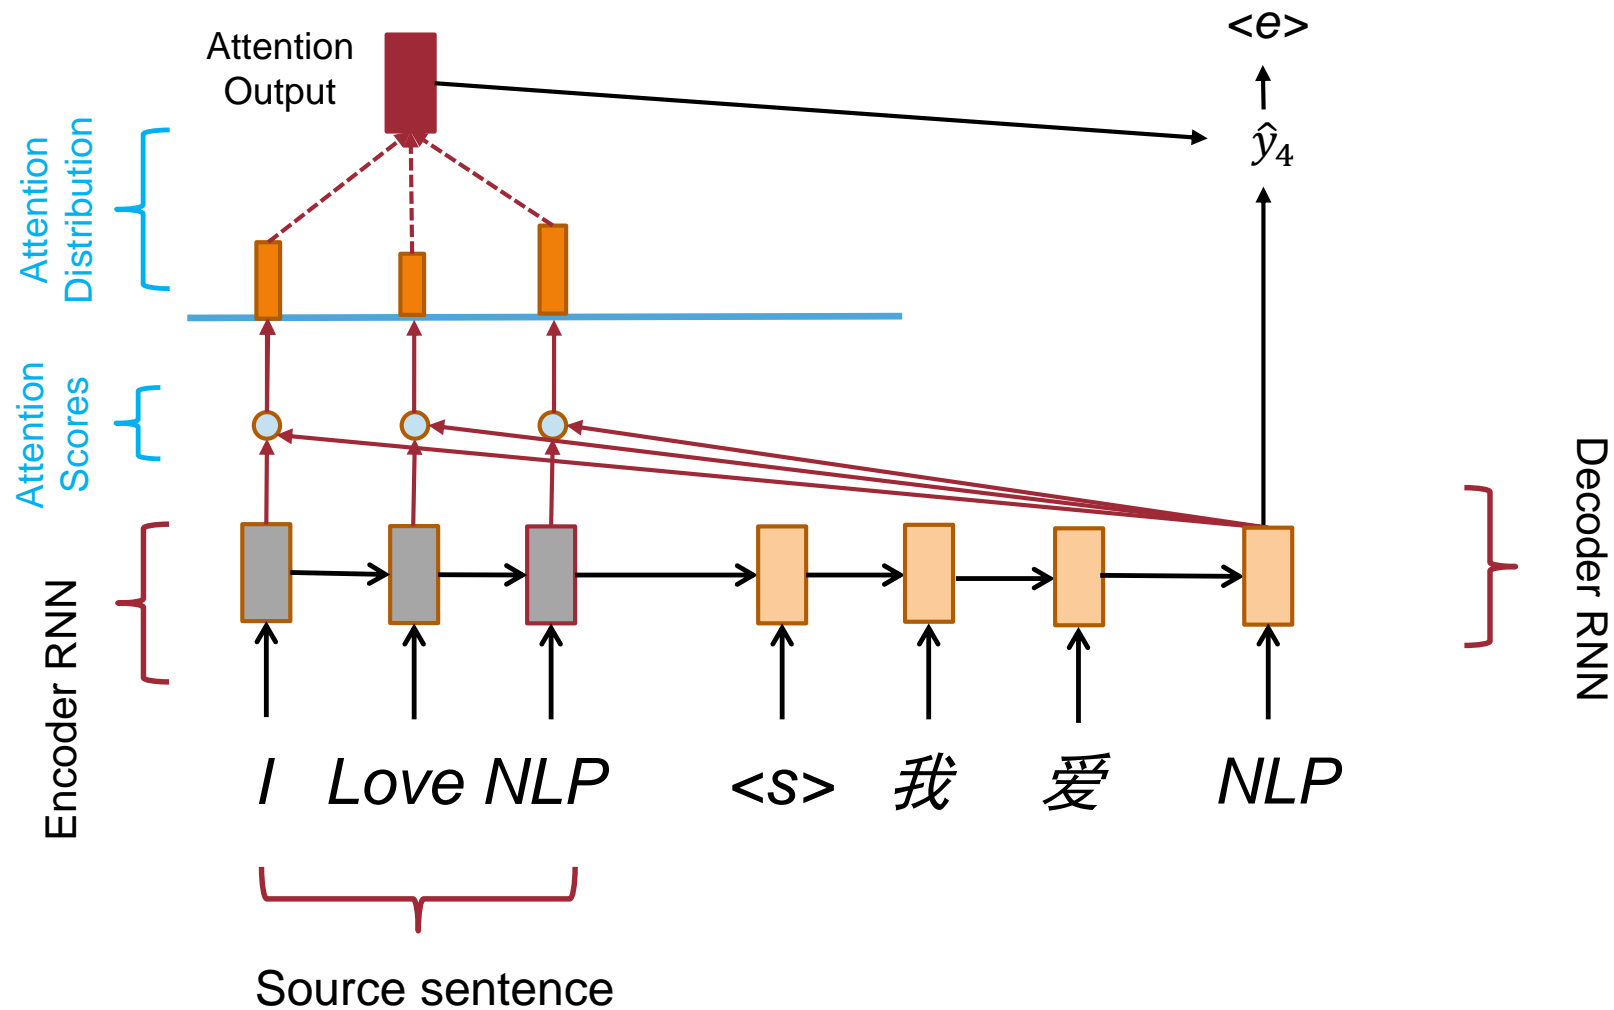
Seq2seq with attention



Seq2seq with attention




Seq2seq with attention



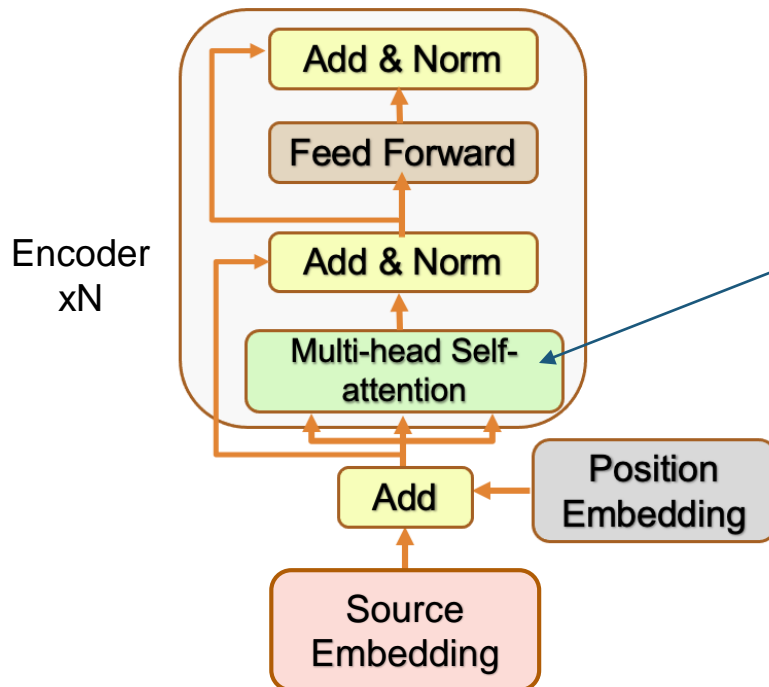
Attention is great

- ▶ Attention solves **the bottleneck problem**
 - ▶ It allows decoder to look directly at source, bypassing bottleneck
- ▶ Attention helps with **the vanishing gradient problem**
 - ▶ Provides shortcut to faraway states
- ▶ Attention provides more “**human-like**” model of seq2seq
 - ▶ You can look back at the source sentence while translating, rather than needing to remember it all
- ▶ Attention provides some **interpretability**
 - ▶ Attention distribution reveals what the decoder was focusing on
 - ▶ = (soft) **source-target sentence alignment**.



Replace RNN with Transformer

- ▶ We have introduced transformer **encoder** in LM.

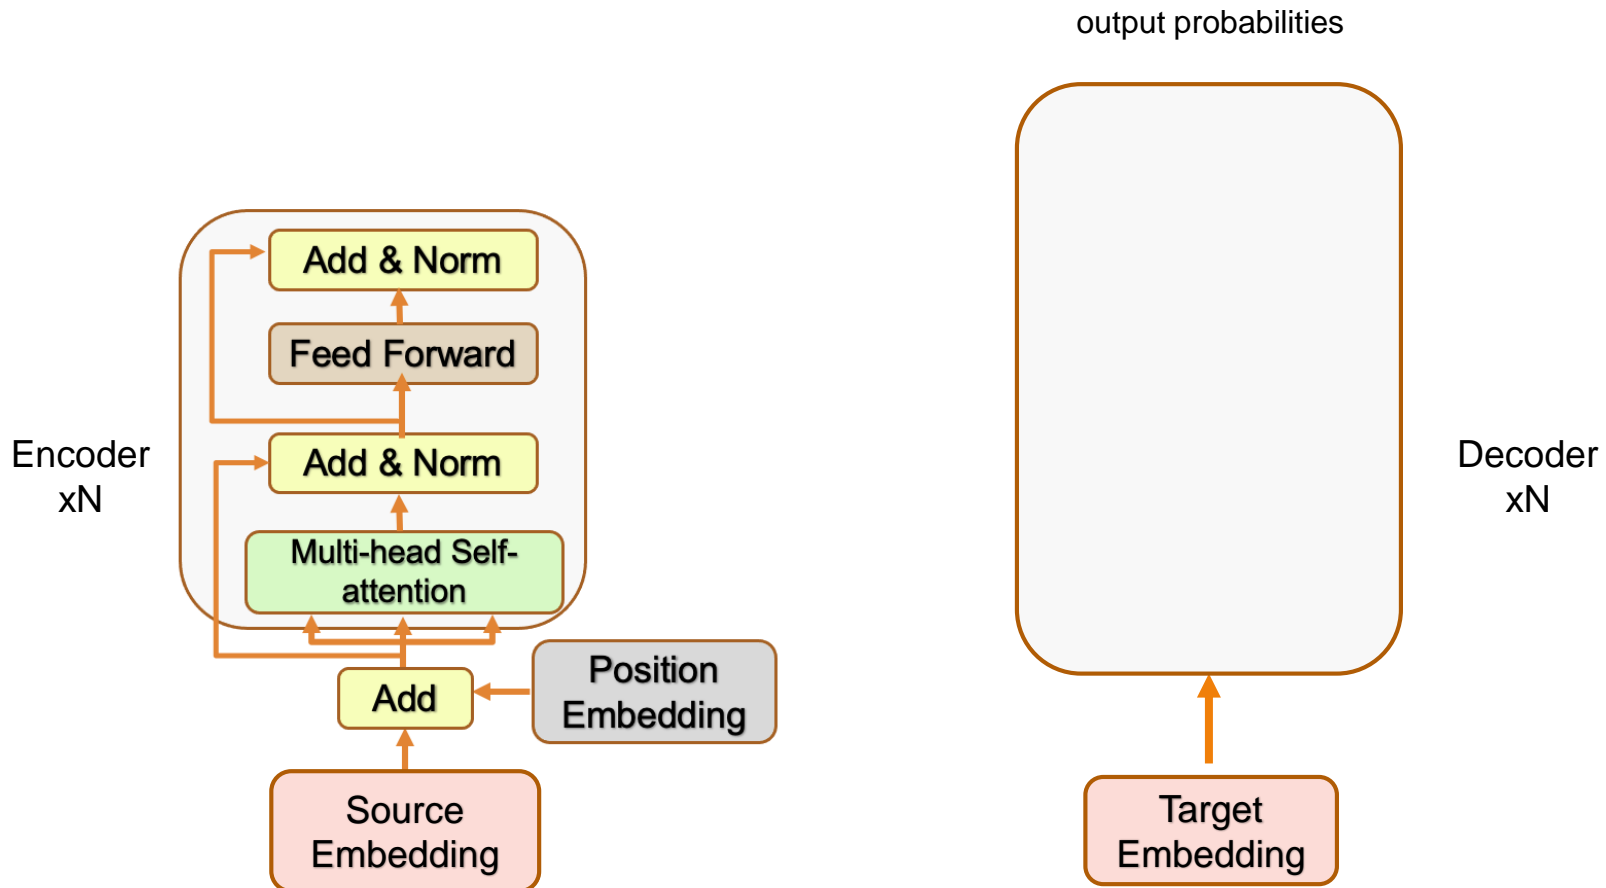


No mask, because we are not doing LM and are allowed to see the complete sentence.



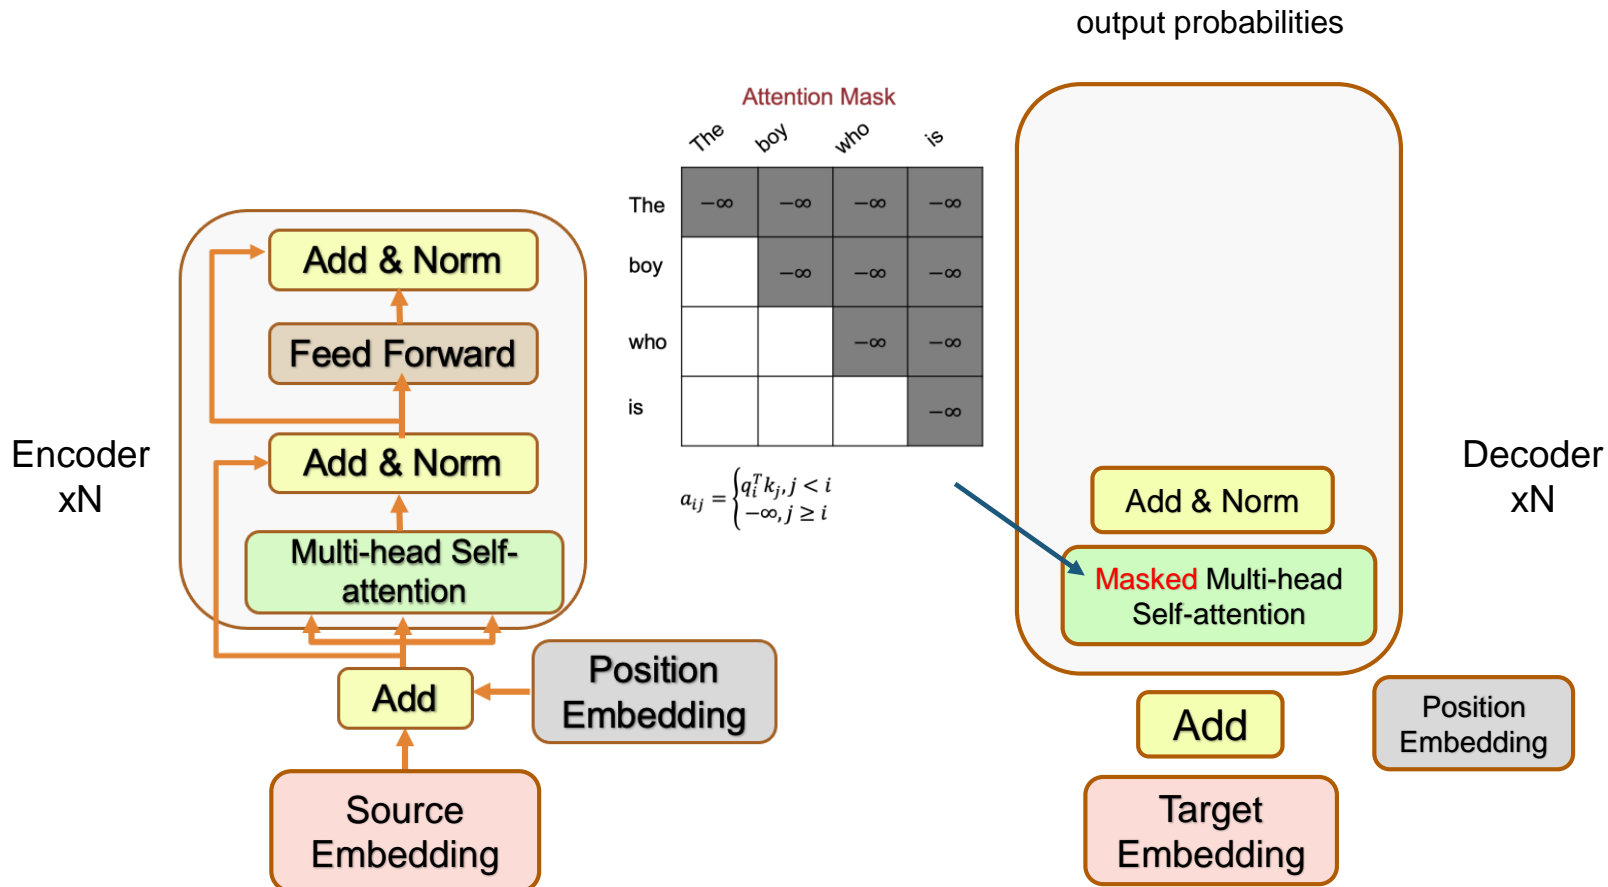
Replace RNN with Transformer

- What about the **decoder**?



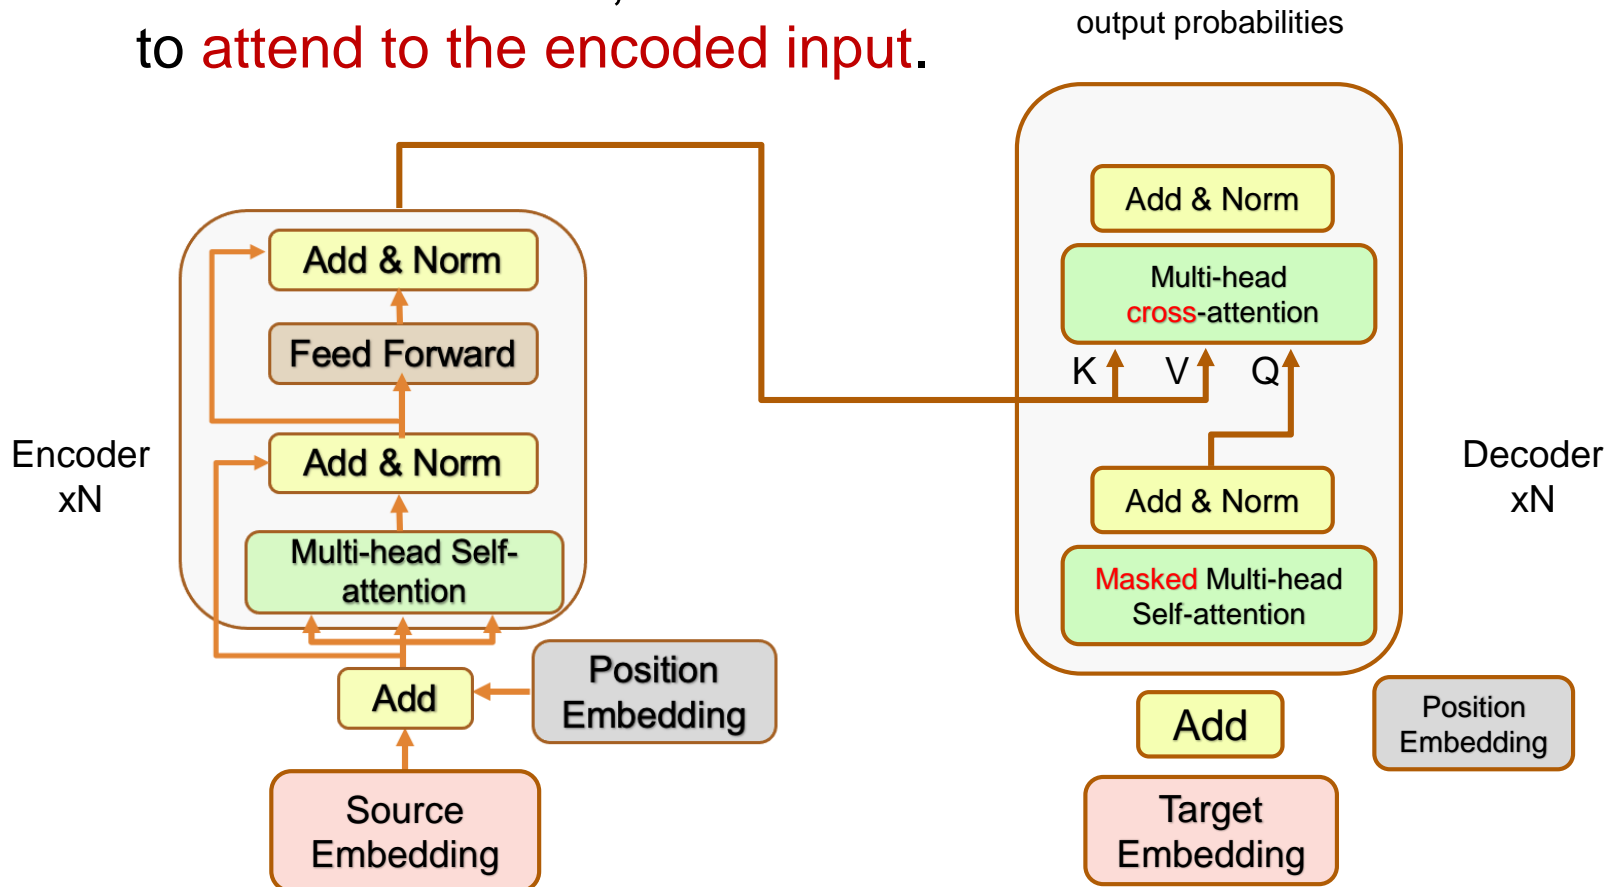
Seq2seq with Transformer

- ▶ We ignore the arrows in the decoder for simplicity.



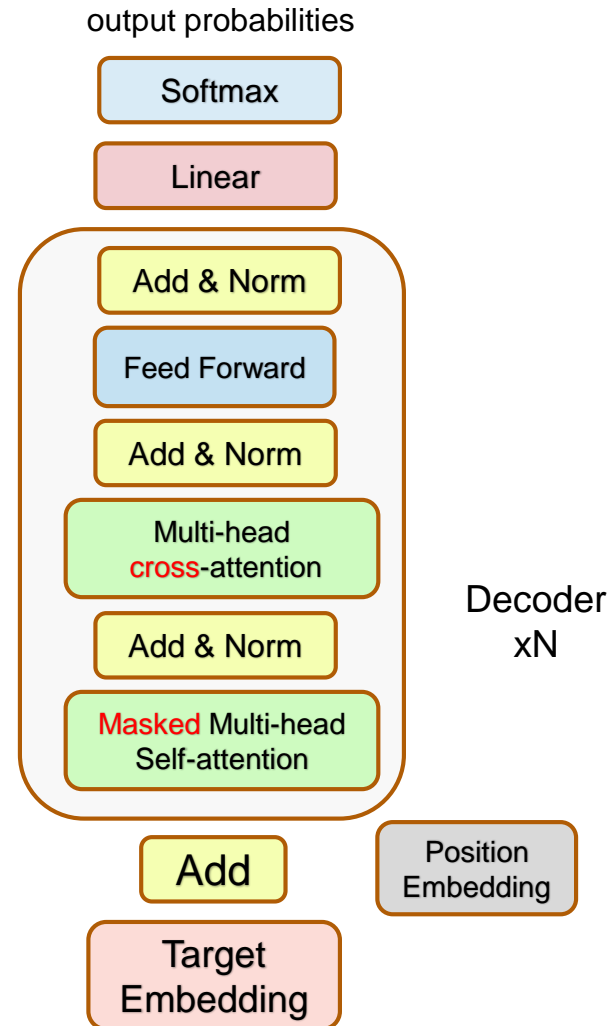
Seq2seq with Transformer

- ▶ In self-attention, K/Q/V come from the same sequence.
- ▶ But in the decoder, we also need to **attend to the encoded input**.

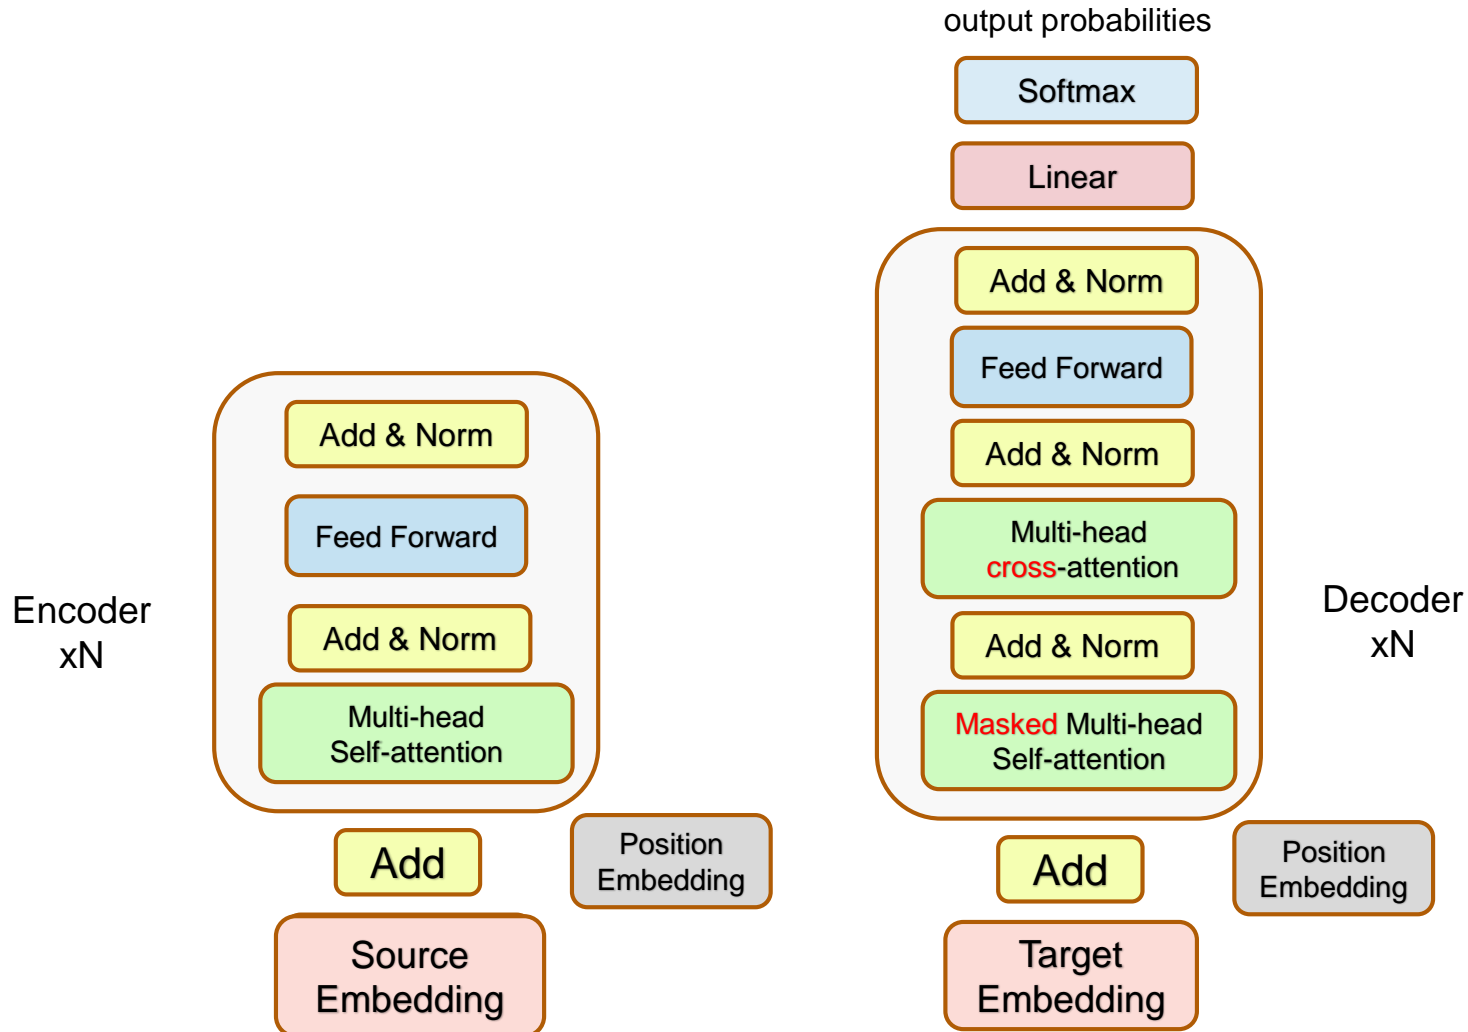


Transformer decoder: finishing touches

- ▶ Add a feed forward layer (with residual connections and layer norm)
- ▶ Add a final linear layer to project the embeddings into a much longer vector of length vocab size (logits)
- ▶ Add a final softmax to generate a probability distribution of next word



Recap of the Transformer Architecture





Learning and Decoding



Learning of seq2seq

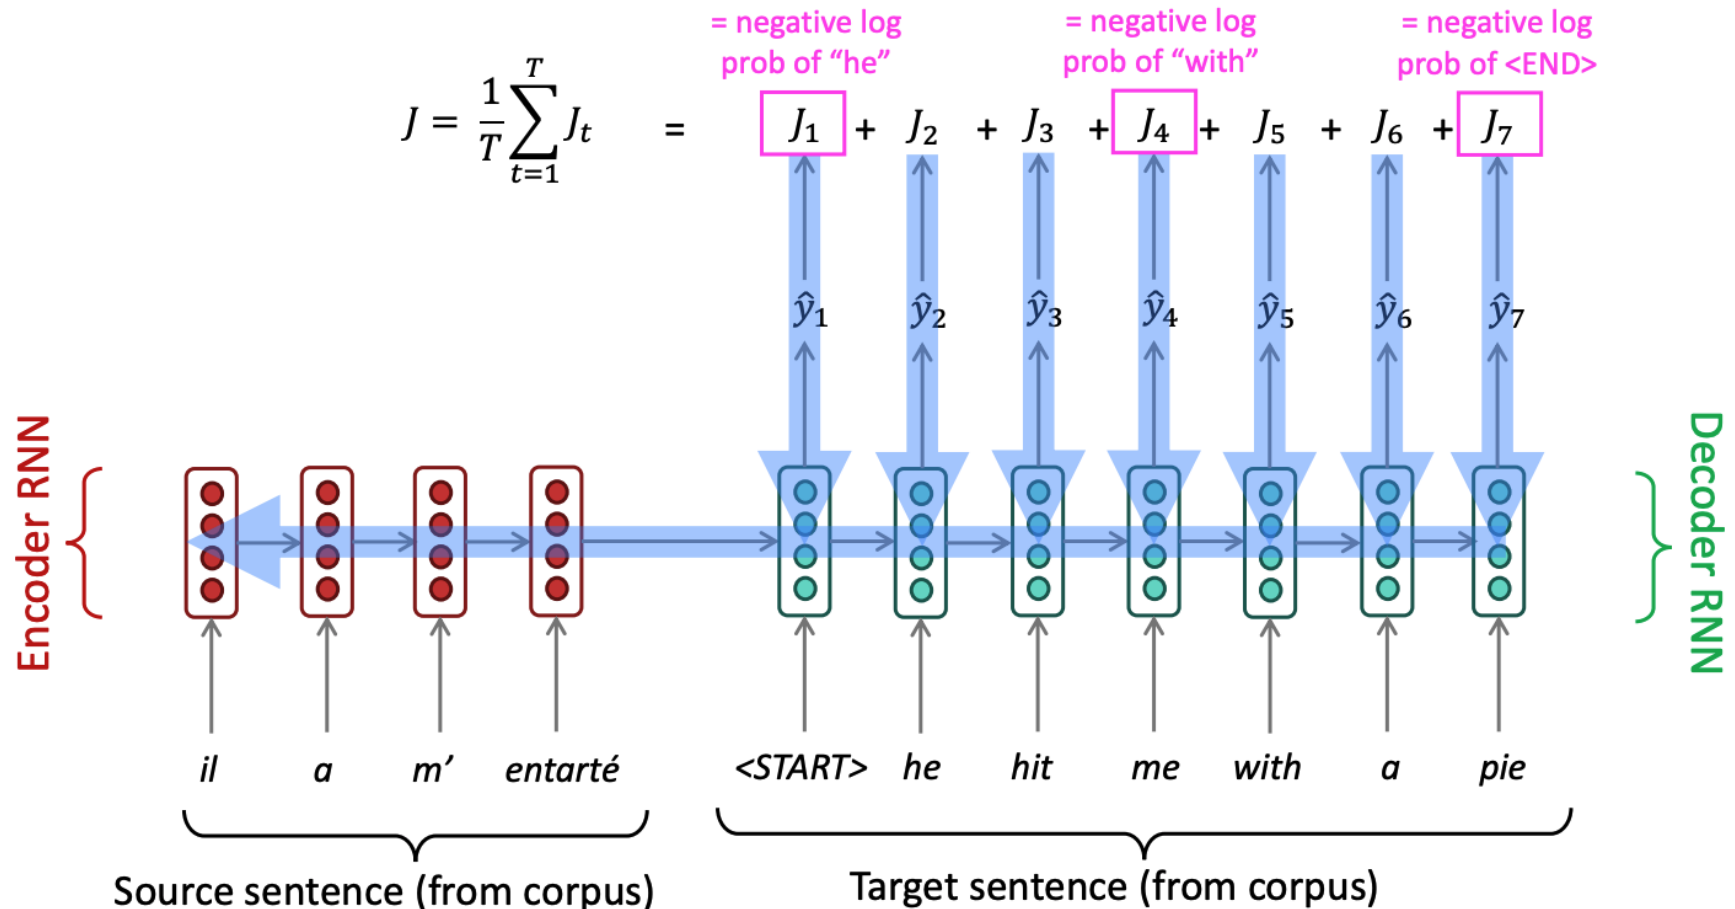
- ▶ The sequence-to-sequence model is an example of a **conditional language model**
 - ▶ “Language model” because the decoder is predicting each word of the target sentence y from left to right
 - ▶ “Conditional” because its predictions are conditioned on the source sentence x
- ▶ Seq2seq directly calculates $P(y|x)$:

$$P(y_{1:T}|x) = P(y_1|x)P(y_2|y_1, x)P(y_3|y_1, y_2, x) \dots P(y_T|y_1, \dots y_{T-1}, x)$$

- ▶ Learning
 - ▶ Training data: a parallel corpus (paired input-output sequences)
 - ▶ Objective: maximizing the conditional likelihood




Learning of seq2seq



Seq2seq is optimized as a **single system**. Backpropagation operates "*end-to-end*".

Greedy decoding

- ▶ We saw how to generate (or “decode”) the target sentence by taking argmax on each step of the decoder
 - ▶ Predict the most probable word at each step



Problems with greedy decoding

- ▶ Greedy decoding has no way to undo decisions!
 - ▶ Input: *il a m'entarté* (*he hit me with a pie*)
 - ▶ → *he* _____
 - ▶ → *he hit* _____
 - ▶ → *he hit a* _____ (*whoops! no going back now...*)
- ▶ How to fix this?



Beam search decoding

- ▶ At each step, keep track of the ***k* best** partial translations (i.e., hypotheses)
 - ▶ *k* is the **beam size** (e.g., 5 to 10 in NMT)
- ▶ A hypothesis y_1, \dots, y_t has a **score** which is its log probability:

$$\text{score}(y_1, \dots, y_t) = \log P_{LM}(y_1, \dots, y_t | x) = \sum_{i=1}^t \log P_{LM}(y_i | y_1, \dots, y_{i-1}, x)$$

- ▶ Beam search is not guaranteed to find optimal solution
- ▶ It is even not guaranteed to outperform greedy decoding
 - ▶ But it is often better



Beam search decoding: example

Beam size = $k = 2$.

Blue numbers = $\text{score}(y_1, \dots, y_t) = \sum_{i=1}^t \log P_{LM}(y_i | y_1, \dots, y_{i-1}, x)$

<start>

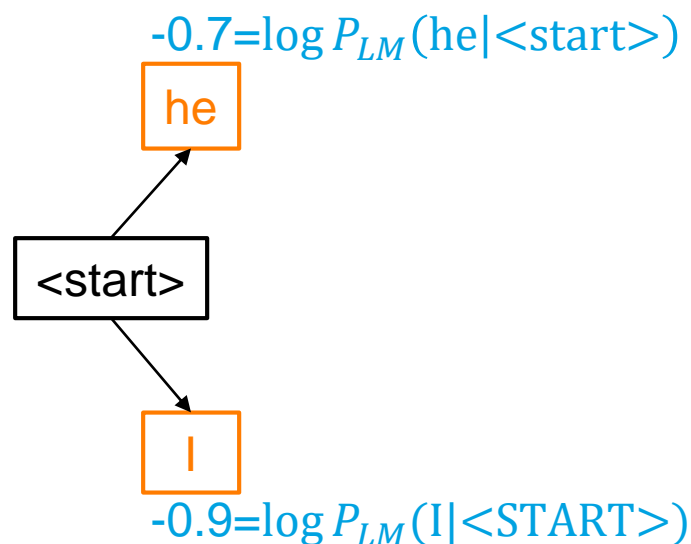
Calculate prob dist of next word



Beam search decoding: example

Beam size = $k = 2$.

Blue numbers = $\text{score}(y_1, \dots, y_t) = \sum_{i=1}^t \log P_{LM}(y_i | y_1, \dots, y_{i-1}, x)$



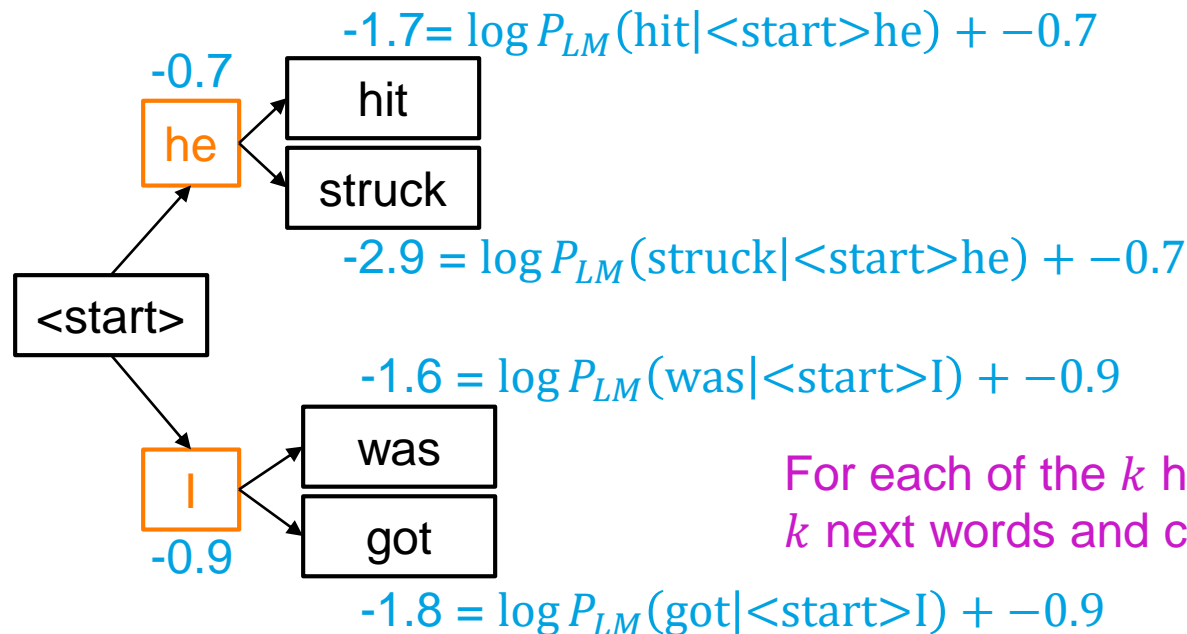
Take top k words



Beam search decoding: example

Beam size = $k = 2$.

Blue numbers = $\text{score}(y_1, \dots, y_t) = \sum_{i=1}^t \log P_{LM}(y_i | y_1, \dots, y_{i-1}, x)$



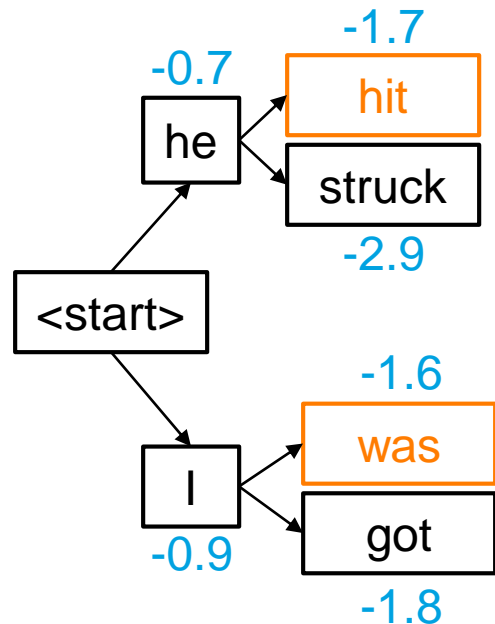
For each of the k hypothesis, find top k next words and calculate scores



Beam search decoding: example

Beam size = $k = 2$.

Blue numbers = $\text{score}(y_1, \dots, y_t) = \sum_{i=1}^t \log P_{LM}(y_i | y_1, \dots, y_{i-1}, x)$



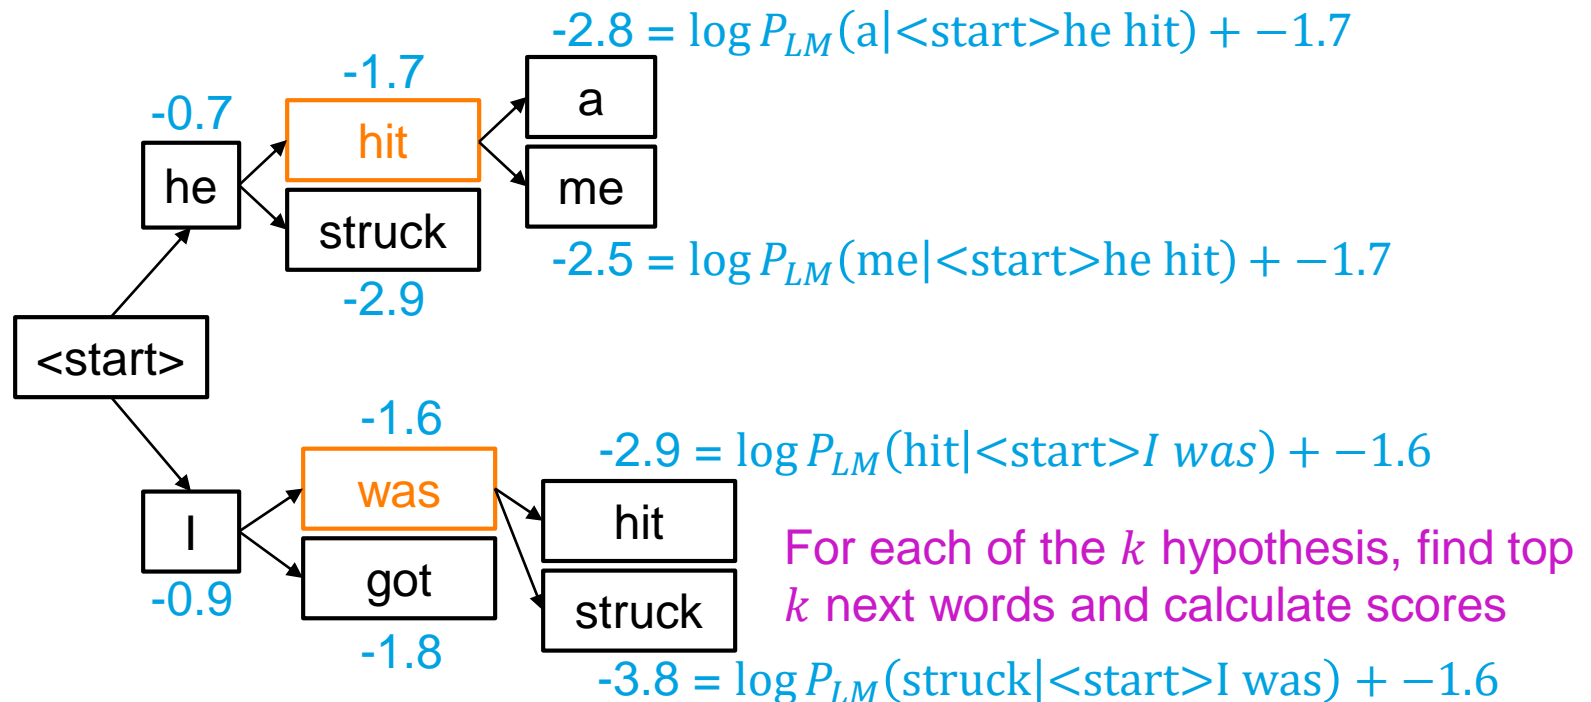
Of these k^2 hypotheses,
just keep k with highest scores



Beam search decoding: example

Beam size = $k = 2$.

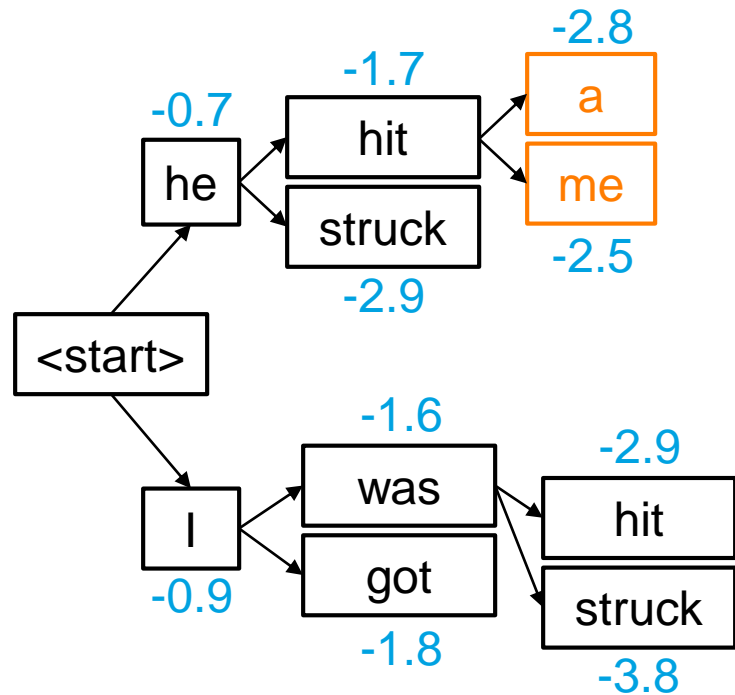
Blue numbers = $\text{score}(y_1, \dots, y_t) = \sum_{i=1}^t \log P_{LM}(y_i | y_1, \dots, y_{i-1}, x)$



Beam search decoding: example

Beam size = $k = 2$.

Blue numbers = $\text{score}(y_1, \dots, y_t) = \sum_{i=1}^t \log P_{LM}(y_i | y_1, \dots, y_{i-1}, x)$

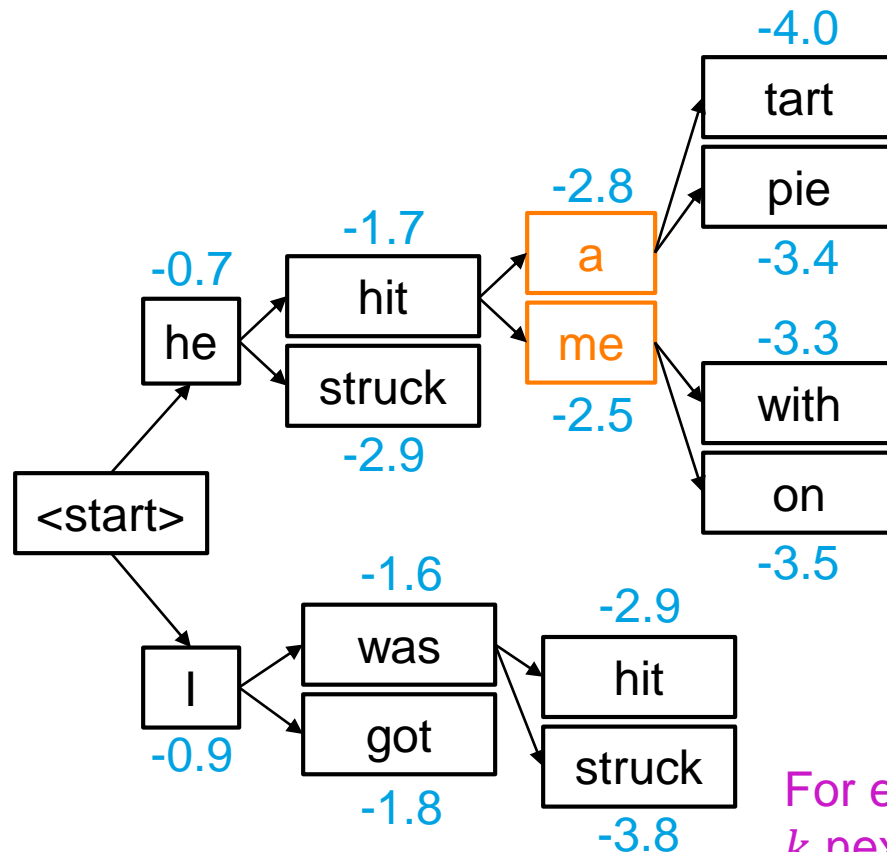


Of these k^2 hypotheses,
just keep k with highest scores

Beam search decoding: example

Beam size = $k = 2$.

Blue numbers = $\text{score}(y_1, \dots, y_t) = \sum_{i=1}^t \log P_{LM}(y_i | y_1, \dots, y_{i-1}, x)$

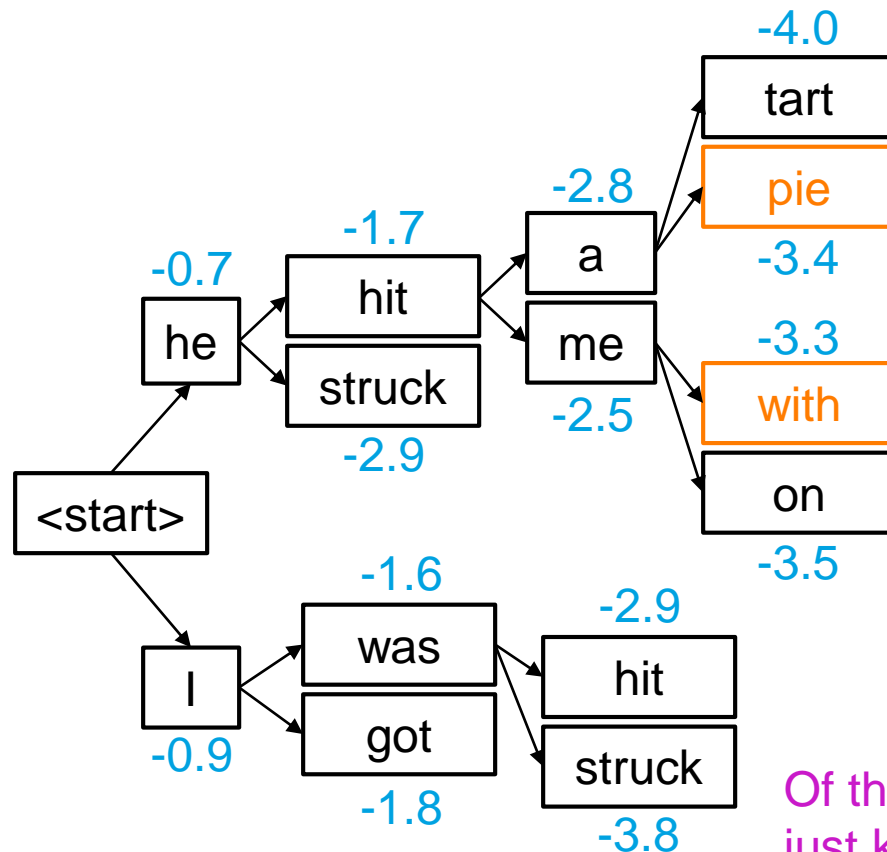


For each of the k hypothesis, find top k next words and calculate scores

Beam search decoding: example

Beam size = $k = 2$.

Blue numbers = $\text{score}(y_1, \dots, y_t) = \sum_{i=1}^t \log P_{LM}(y_i | y_1, \dots, y_{i-1}, x)$

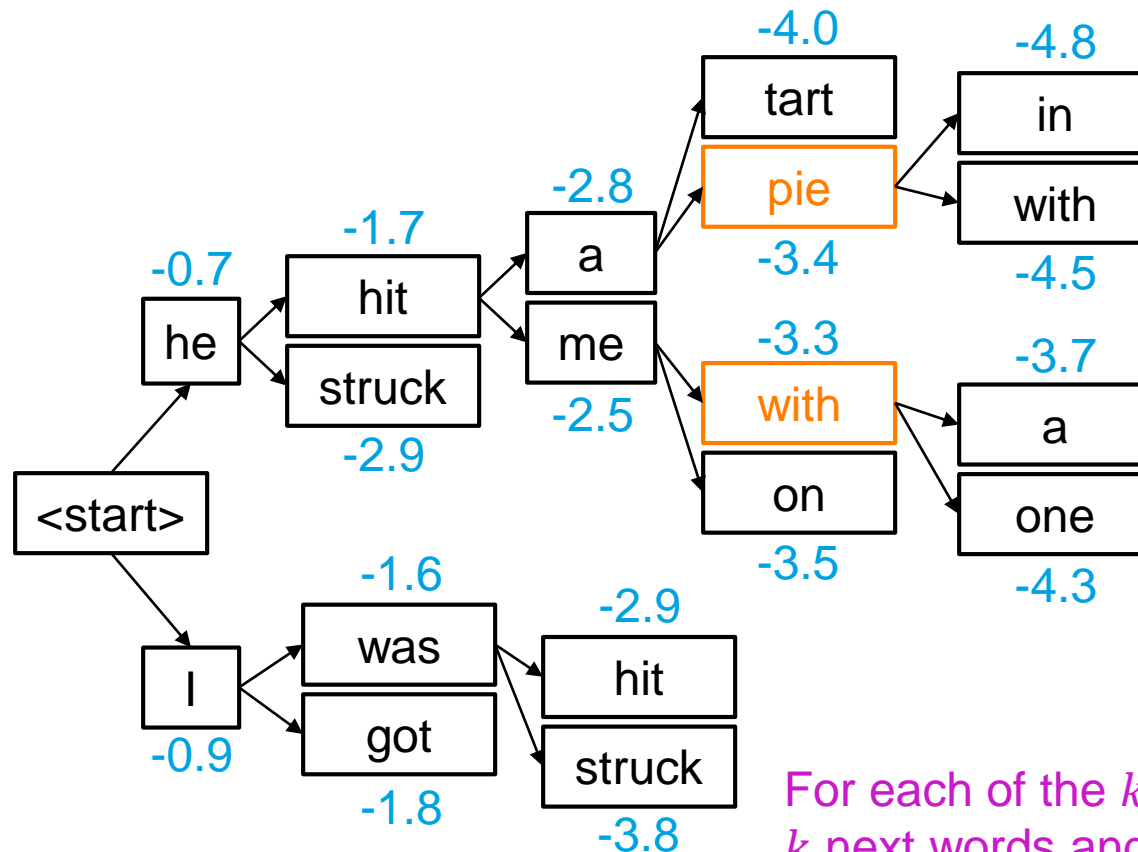


Of these k^2 hypotheses,
just keep k with highest scores

Beam search decoding: example

Beam size = $k = 2$.

Blue numbers = $\text{score}(y_1, \dots, y_t) = \sum_{i=1}^t \log P_{LM}(y_i | y_1, \dots, y_{i-1}, x)$

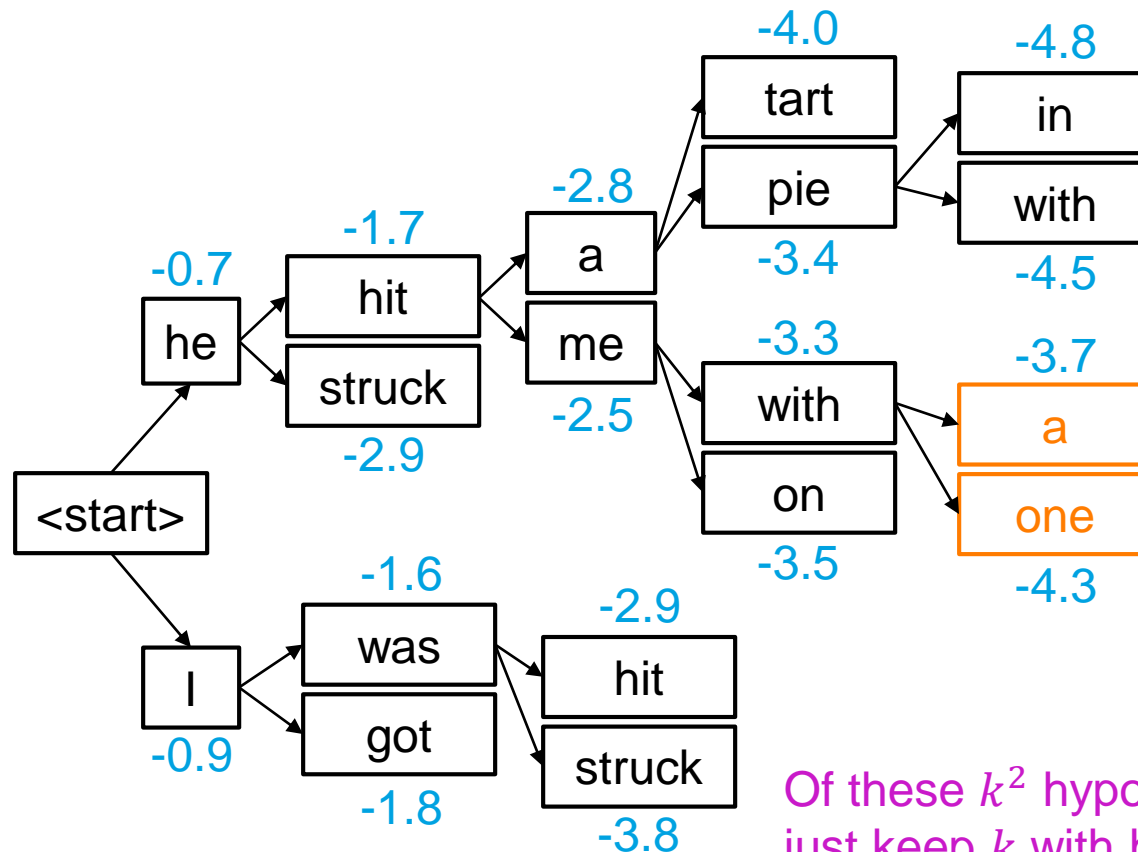


For each of the k hypothesis, find top k next words and calculate scores

Beam search decoding: example

Beam size = $k = 2$.

Blue numbers = $\text{score}(y_1, \dots, y_t) = \sum_{i=1}^t \log P_{LM}(y_i | y_1, \dots, y_{i-1}, x)$

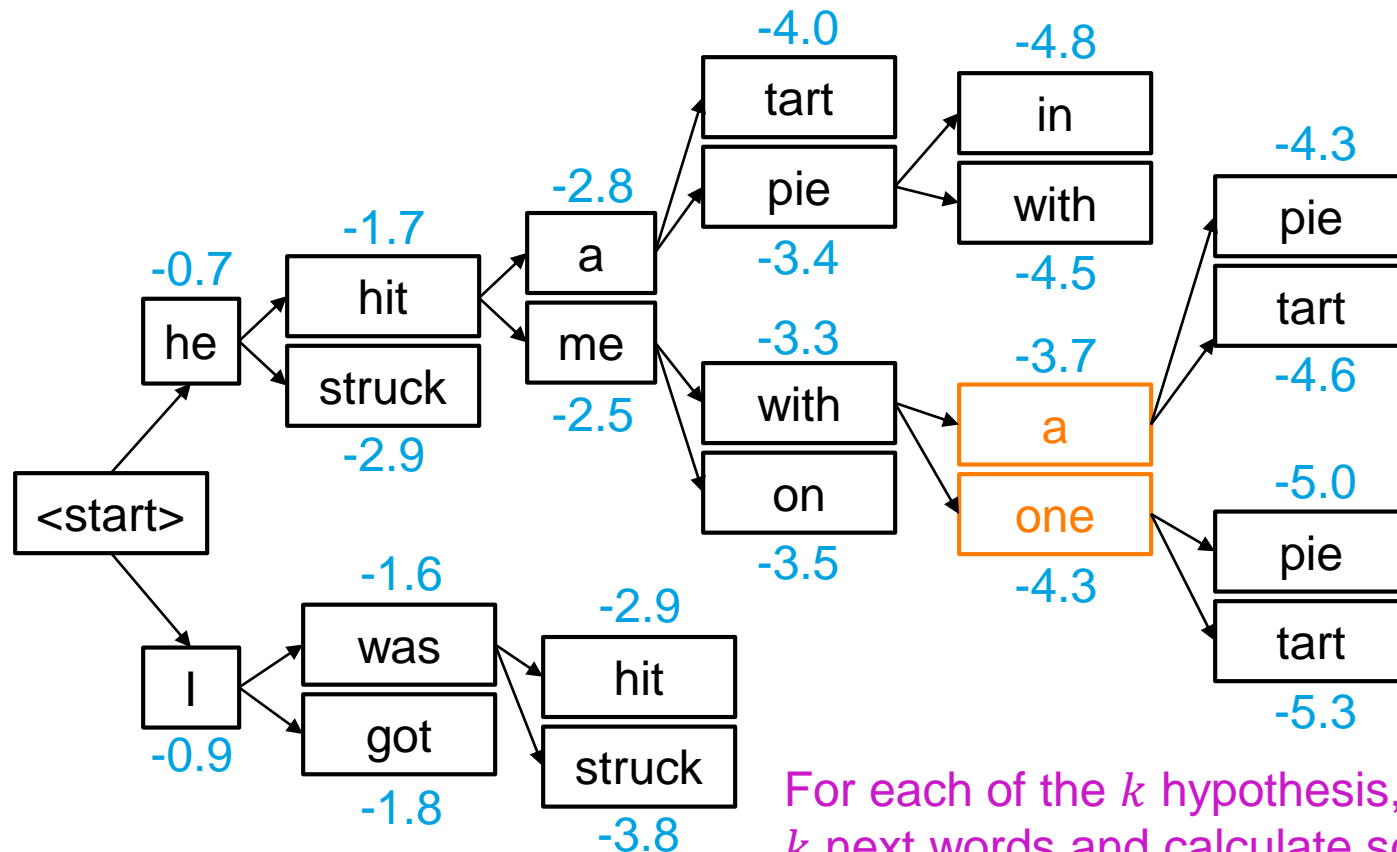


Of these k^2 hypotheses,
just keep k with highest scores

Beam search decoding: example

Beam size = $k = 2$.

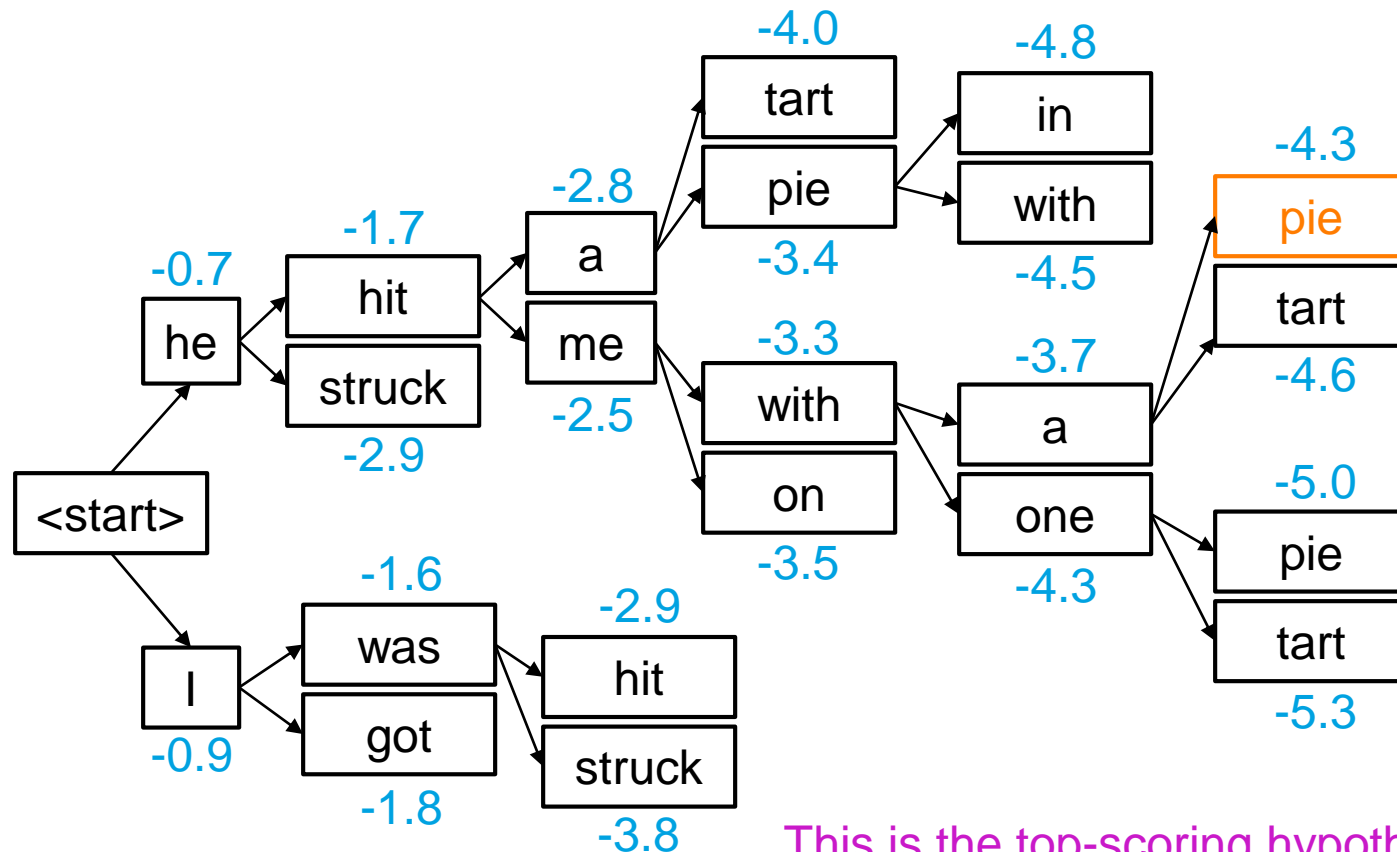
Blue numbers = $\text{score}(y_1, \dots, y_t) = \sum_{i=1}^t \log P_{LM}(y_i | y_1, \dots, y_{i-1}, x)$



Beam search decoding: example

Beam size = $k = 2$.

Blue numbers = $\text{score}(y_1, \dots, y_t) = \sum_{i=1}^t \log P_{LM}(y_i | y_1, \dots, y_{i-1}, x)$

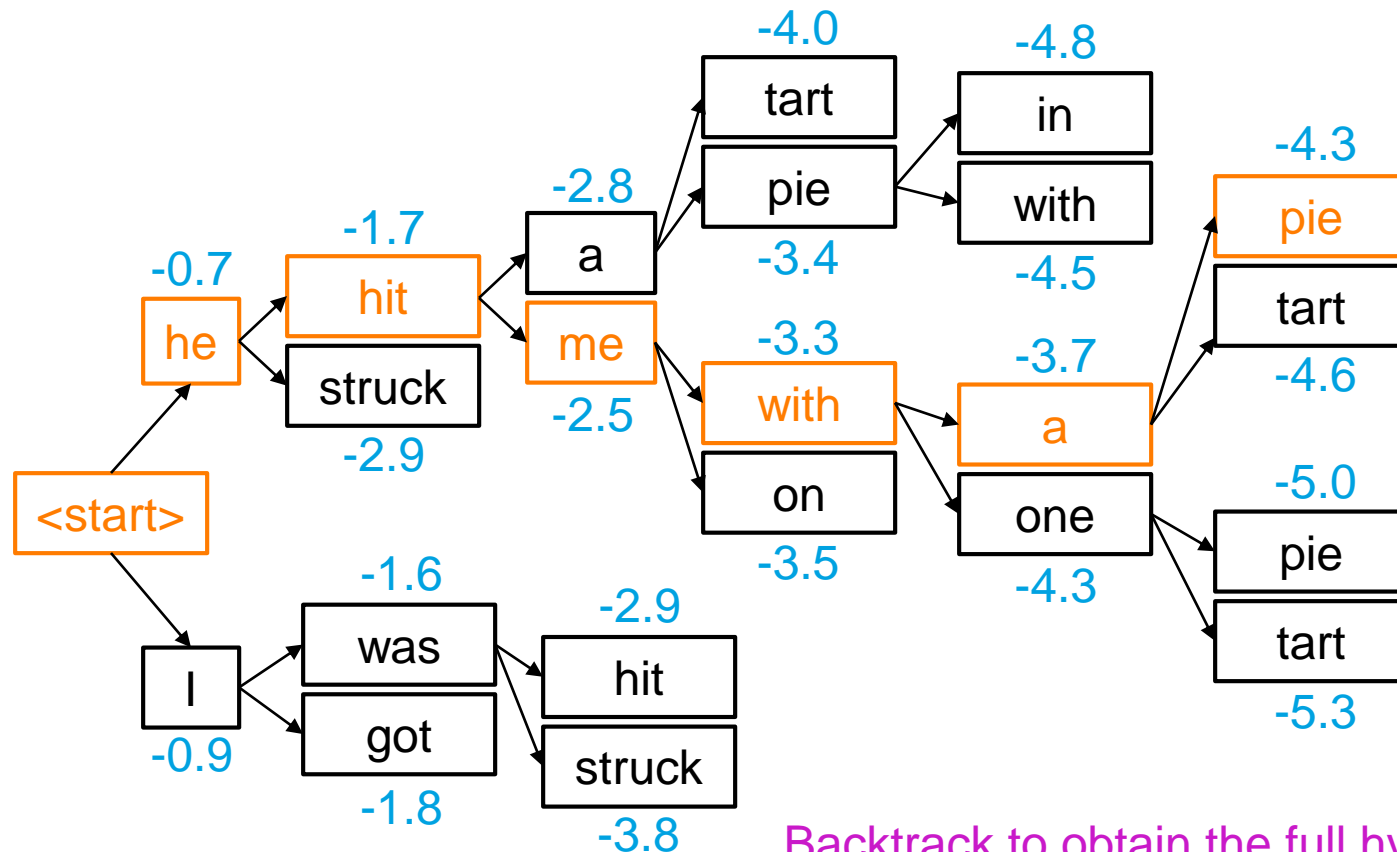


This is the top-scoring hypothesis!

Beam search decoding: example

Beam size = $k = 2$.

Blue numbers = $\text{score}(y_1, \dots, y_t) = \sum_{i=1}^t \log P_{LM}(y_i | y_1, \dots, y_{i-1}, x)$



Backtrack to obtain the full hypothesis

Beam search decoding: stopping criterion

- ▶ In greedy decoding, usually we decode until the model produces an **<END>** token
 - ▶ For example, “<START> *he hit me with a pie* <END>”
- ▶ In beam search decoding, different hypotheses may produce <END> tokens on **different timesteps**
 - ▶ When a hypothesis produces <END>, that hypothesis is **complete**.
 - ▶ Place it aside and continue exploring other hypotheses via beam search.
- ▶ Usually we continue beam search until:
 - ▶ We reach timestep T , or
 - ▶ We have at least n completed hypotheses

T and n are some pre-defined cutoffs



Beam search decoding: finishing up

- ▶ With a list of completed hypotheses, how do we select the best one?
- ▶ Selecting one with the highest score?

$$\text{score}(y_1, \dots, y_t) = \log P_{LM}(y_1, \dots, y_t | x) = \sum_{i=1}^t \log P_{LM}(y_i | y_1, \dots, y_{i-1}, x)$$

- ▶ Problem with this?
- ▶ Longer hypotheses have lower scores
- ▶ Fix: Normalize by length. Use this to select the top one:

$$\frac{1}{t} \sum_{i=1}^t \log P_{LM}(y_i | y_1, \dots, y_{i-1}, x)$$



Decoding

- ▶ Other desiderata

- ▶ Reduce repetition

- ▶ Penalize generation of already-seen tokens

- ▶ Can also be incorporated into training (unlikelihood objective)

- ▶ Prevent attending to the same words

- ▶ Encourage diversity for certain tasks (e.g., dialog)

- ▶ Sample, instead of being greedy

- ▶ Avoid hallucination (generating things not in the input) for translation, summarization, etc.

- ▶ Multiple causes and solutions





Extensions



Pointer Net / Copy Mechanism

- ▶ In some cases, we may wish to copy part of the input to the output
 - ▶ Ex: named entities, numbers

Source sentence: Federer suffered a 1-3 defeat to Nadal

Target sentence: Federer a subi une défaite 1-3 contre Nadal



Pointer Net / Copy Mechanism

- ▶ In some cases, we may wish to copy part of the input to the output
 - ▶ Ex: semantic role labeling as seq2seq

Source sentence:

Tolkien's epic novel The Lord of the Rings was published in 1954-1955 ...

Target output:

Tolkien's epic novel [The Lord of the Rings | *subject*]
[was published | *predicate*] [in 1954-1955 | *temporal*]

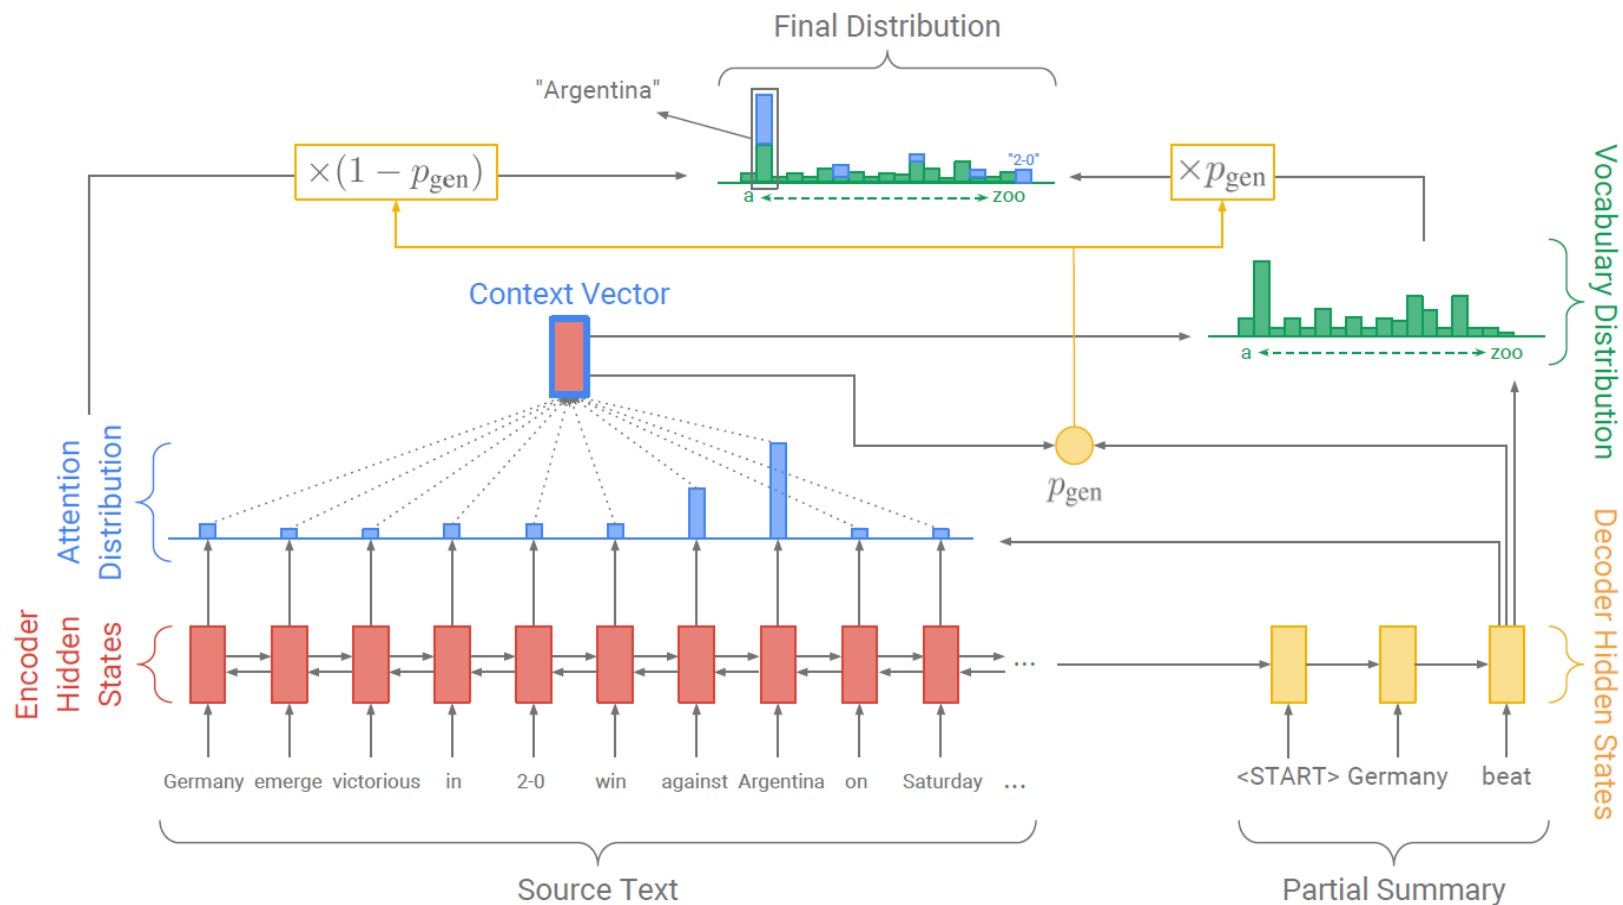


Pointer Net / Copy Mechanism

- ▶ At each decoding step, calculate a Bernoulli distribution over **generation vs. copying** and sample from it
 - ▶ Generation: the old way of predicting a token
 - ▶ Copying: predict a distribution over the input tokens (i.e., a pointer)
 - ▶ This is similar to computing the attention distribution



Pointer Net / Copy Mechanism



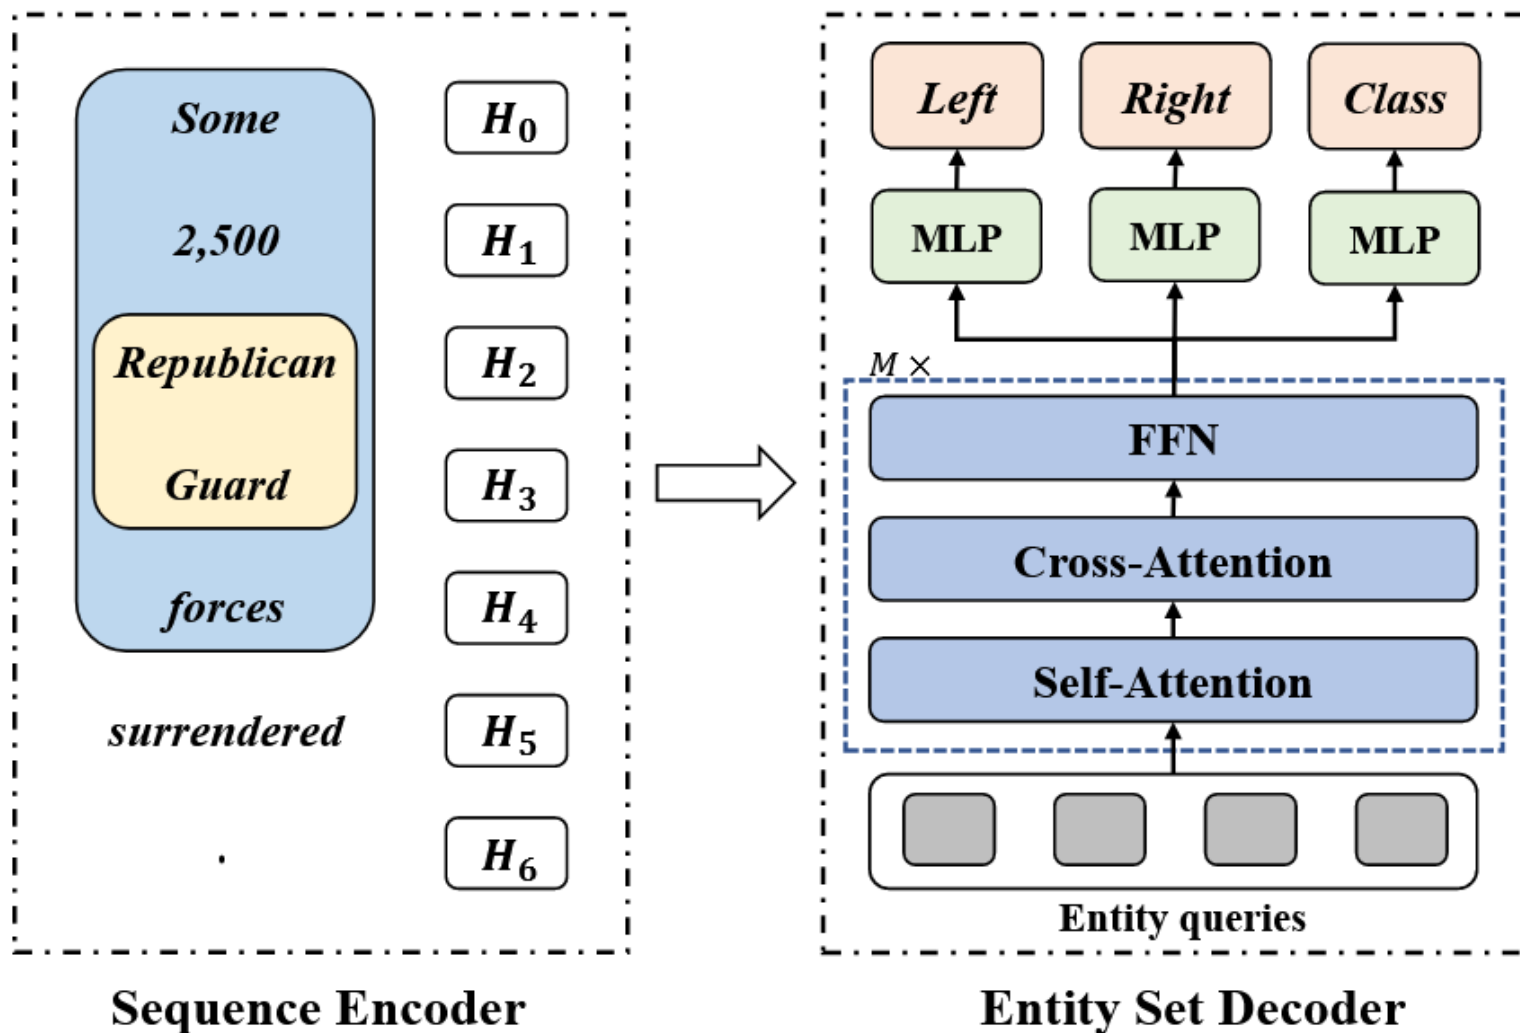
Extensions of Seq2Seq

▶ Seq2Set

- ▶ Decoding to a set (e.g., named entities)
- ▶ Still use a Transformer decoder, but the input is a fixed number of “queries”
 - ▶ Query vectors are learnable model parameters
 - ▶ No position embedding, no mask
 - ▶ Simultaneous decoding across all the positions
 - ▶ An output can be null, allowing variable set sizes



Extensions of Seq2Seq



Extensions of Seq2Seq

- ▶ X2Seq (requiring an X-encoder)
 - ▶ Image captioning
 - ▶ Visual question answering
 - ▶ Structured data (e.g., table) to text
 - ▶ ...
- ▶ Null2Seq
 - ▶ Language modeling!
 - ▶ (Unconditional) text generation





Summary



Sequence to Sequence

- ▶ Many applications
 - ▶ MT, paraphrase, summarization, ...
- ▶ Methods: encoder-decoder
 - ▶ Recurrent neural network
 - ▶ Attention
 - ▶ Transformer: cross-attention
- ▶ Learning
 - ▶ Maximizing conditional likelihood on a parallel corpus
- ▶ Decoding
 - ▶ Greedy, beam-search
- ▶ Extensions
 - ▶ Pointer Net / Copy Mechanism
 - ▶ Seq2Set, X2Seq, Null2Seq

